

Plant_Health_Model

November 28, 2021

1 Visual Deep Learning: Plant Health Model

1.0.1 Gabriel Beutl & Alexander Schwab

This notebook describes the creation of a model to evaluate the health of plants. The model was trained with leaves from an apple tree. With this, an application should be able to identify if a tree is unhealthy, by providing a image form a leave of the tree.

The idea to create this model is based on a project that was started last semester. The project GreenThumb is a plant health tracking app, that allows the user to keep track of the health of a plant. With the addition of a trained ai model, the application could be improved a lot in its functionality.

All the steps for the creation of this model are described point by point in this notebook.

1.0.2 Steps to create this Machine Learning Model:

1. Setup
2. Load image data and apply data augmentation
3. Define model
4. Display model structure
5. Training model
6. Write history and plot graphs
7. Evaluate Model
8. Predict with model and display ROC
9. Save trained model
10. Load trained model
11. Test trained model

1.0.3 Experiments:

The notebook tries different experiments for the creation of the best model. As such, steps 3 - 8 are repeated three times with different models.

1.0.4 1. Setup

```
[1]: import numpy as np
import tensorflow as tf
from tensorflow.keras import metrics
from sklearn import metrics as skmetrics

from sklearn.metrics import confusion_matrix, accuracy_score

import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator

[2]: #Environment Variables
EPOCHS = 50
BATCH_SIZE = 32
IMG_SIZE = (256, 256)
LEARNING_RATE = 0.001
MODEL_NAME = 'plant_health_model.h5'

#Folder for the data
train_data_dir = './plant_images/train'
valid_data_dir = './plant_images/valid'
test_data_dir = './plant_images/test'

#Different Metrics that get printed in the history
METRICS = [
    metrics.TruePositives(name='tp'),
    metrics.FalsePositives(name='fp'),
    metrics.TrueNegatives(name='tn'),
    metrics.FalseNegatives(name='fn'),
    metrics.CategoricalAccuracy(name='accuracy'),
    metrics.Precision(name='precision'),
    metrics.AUC(name='auc'),
    metrics.MeanAbsoluteError(name='mae'),
    metrics.MeanSquaredError(name='mse')
]
```

1.0.5 2. Load Image Data and apply data augmentation and preprocessing

The structure of the folder is the following: -root -plant_images —train —Apple_healthy —Apple_unhealthy —valid —Apple_healthy —Apple_unhealthy —test —Apple_healthy —Apple_unhealthy

The train folder includes around 1400 images for the training. The valid folder includes around 700 images to validate the trained model. The test folder includes again around 120 images to test the model.

Augmentation: For the augmentation a ImageDataGenerator is used. With this the images for the training and validation are rescaled, sheared, rotated, flipped and zoomed. The test images are only rescaled.

```
[3]: #Create augmented training data
train_datagen = ImageDataGenerator(rescale=(1. / 255),
                                   shear_range=0.1,
                                   zoom_range=0.1,
                                   rotation_range=10,
                                   horizontal_flip=True,
                                   fill_mode='constant',
                                   validation_split=0.2,
                                   cval=0)

#Create augmented training data
train_dataset = train_datagen.flow_from_directory(train_data_dir,
                                                  target_size=IMG_SIZE,
                                                  batch_size=BATCH_SIZE,
                                                  class_mode='categorical',
                                                  subset='training')

#Create augmented validation data
validation_dataset = train_datagen.flow_from_directory(valid_data_dir,
                                                       target_size=IMG_SIZE,
                                                       batch_size=BATCH_SIZE,
                                                       class_mode='categorical',
                                                       subset='validation')

test_datagen = ImageDataGenerator(rescale=(1. / 255))

#Create test data
test_dataset = test_datagen.flow_from_directory(test_data_dir,
                                                target_size=IMG_SIZE,
                                                batch_size=1,
                                                shuffle=False,
                                                class_mode='categorical')
```

Found 1458 images belonging to 2 classes.

Found 126 images belonging to 2 classes.

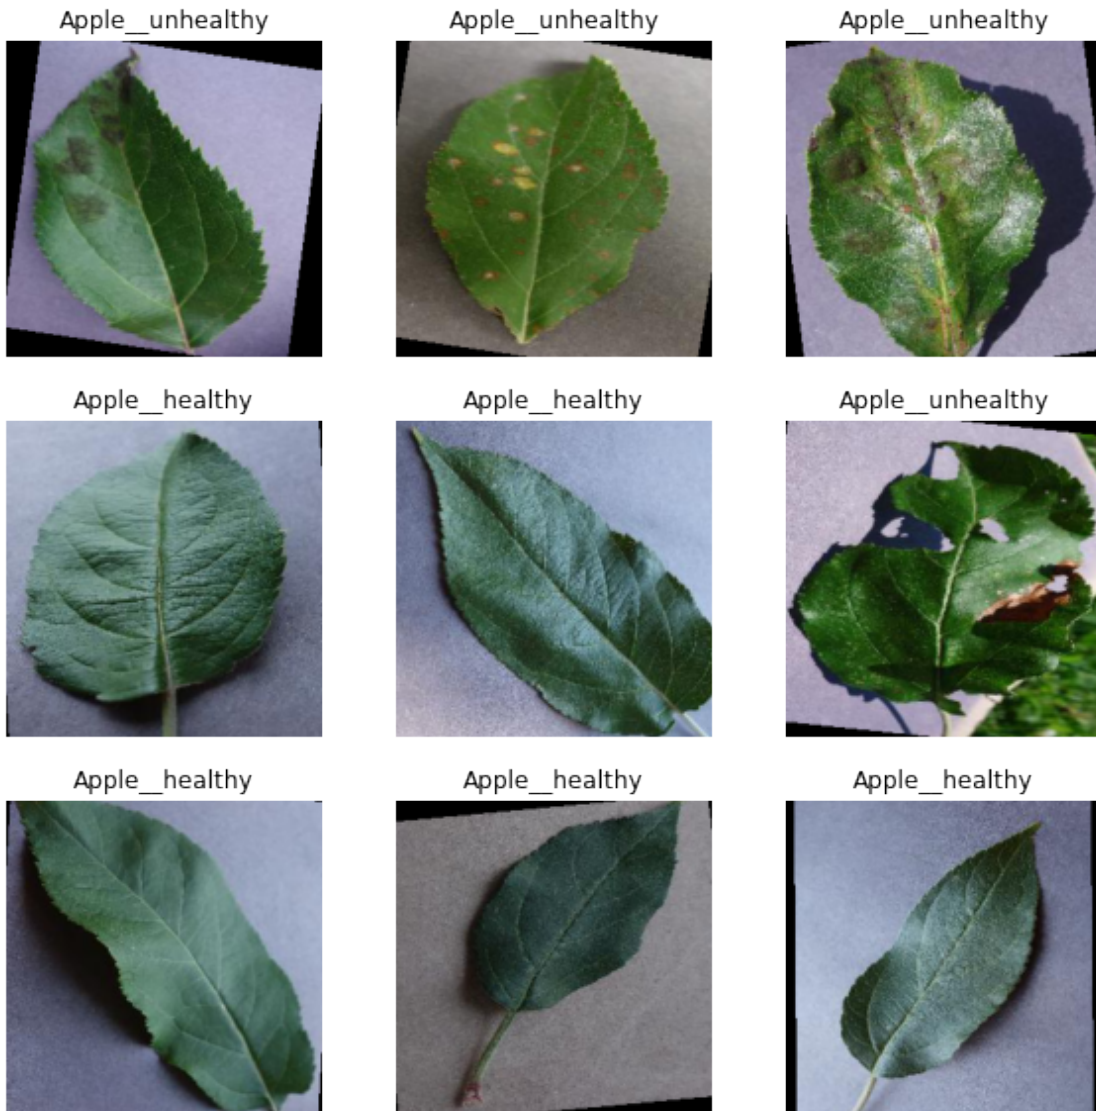
Found 714 images belonging to 2 classes.

```
[4]: # Checks if classes where found and displays them
num_classes = train_dataset.num_classes
class_names = list(train_dataset.class_indices.keys())
print(class_names)
```

```
['Apple__healthy', 'Apple__unhealthy']
```

1.0.6 Display augmented image data

```
[5]: # display the random validation data for testing
x,y = validation_dataset.next()
labels = list(np.argmax(l) for l in y)
plt.figure(figsize=(10, 10))
for i in range(9):
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(x[i])
    plt.title(class_names[labels[i]])
    plt.axis("off")
```



1.1 Here the experiments start with different models

1.1.1 3. Create and compile model 1

This model uses two hidden convolution layers and one dense layer to analyse the images. In between the hidden layers, max-pooling manipulates the values, to create a better learning base. Based on literature the best hidden layer activation is the relu function and for the dense activation softmax. For the loss function, the categorical_crossentropy or the binary_crossentropy (since there are only two classes) was possible for image recognition. For this example, categorical_crossentropy worked better.

```
[6]: model1 = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(128, (3, 3), input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(128, (3, 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```

```
[7]: opt = tf.keras.optimizers.Adam(lr=LEARNING_RATE)
model1.compile(optimizer=opt,
               loss='categorical_crossentropy',
               metrics=METRICS)
```

1.1.2 4. Display model structure 1

```
[8]: model1.summary()
```

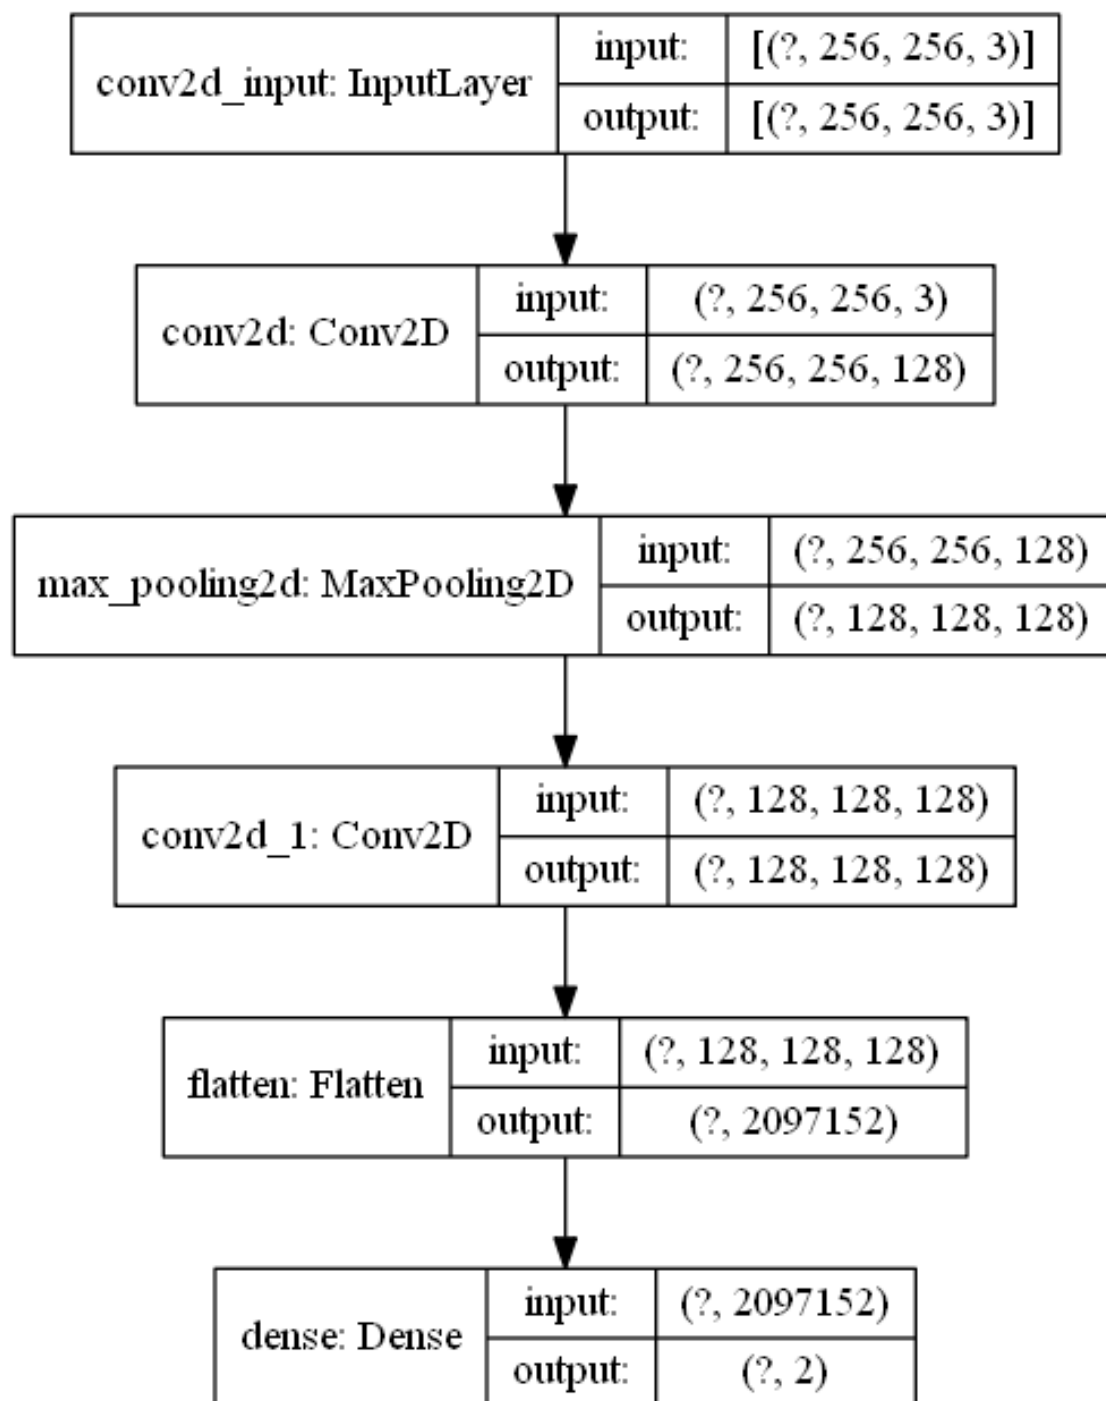
Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|-----------------------|---------|
| conv2d (Conv2D) | (None, 256, 256, 128) | 3456 |
| max_pooling2d (MaxPooling2D) | (None, 128, 128, 128) | 0 |
| conv2d_1 (Conv2D) | (None, 128, 128, 128) | 147456 |
| flatten (Flatten) | (None, 2097152) | 0 |
| dense (Dense) | (None, 2) | 4194306 |

Total params: 4,345,218
Trainable params: 4,345,218
Non-trainable params: 0

```
[9]: tf.keras.utils.plot_model(model1, show_shapes=True)
```

[9]:



1.1.3 5. Training model 1

```
[10]: def train_model(model, model_name):
    #filepath = './models/20210610_plant_' + model_name + '_{epoch:
    ↳02d}-{val_accuracy:.4f}.hdf5'
    #checkpoint = tf.keras.callbacks.ModelCheckpoint(filepath,
    ↳monitor='val_accuracy', verbose=1,
    #
    save_best_only=True,
    ↳mode='max')
    #callbacks_list = [checkpoint]

    history = model.fit(train_dataset,
                        epochs=EPOCHS,
                        validation_data=validation_dataset)
                        #, callbacks=callbacks_list)

    return history

history1 = train_model(model1, "Model_1")
```

Epoch 1/50

46/46 [=====] - 16s 356ms/step - loss: 3.1057 - tp: 956.0000 - fp: 502.0000 - tn: 956.0000 - fn: 502.0000 - accuracy: 0.6557 - precision: 0.6557 - auc: 0.6797 - mae: 0.4453 - mse: 0.2502 - val_loss: 0.4508 - val_tp: 104.0000 - val_fp: 22.0000 - val_tn: 104.0000 - val_fn: 22.0000 - val_accuracy: 0.8254 - val_precision: 0.8254 - val_auc: 0.9048 - val_mae: 0.3332 - val_mse: 0.1428

Epoch 2/50

46/46 [=====] - 16s 338ms/step - loss: 0.3775 - tp: 1210.0000 - fp: 248.0000 - tn: 1210.0000 - fn: 248.0000 - accuracy: 0.8299 - precision: 0.8299 - auc: 0.9110 - mae: 0.2464 - mse: 0.1199 - val_loss: 0.2627 - val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 - val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9607 - val_mae: 0.1745 - val_mse: 0.0766

Epoch 3/50

46/46 [=====] - 16s 337ms/step - loss: 0.3092 - tp: 1278.0000 - fp: 180.0000 - tn: 1278.0000 - fn: 180.0000 - accuracy: 0.8765 - precision: 0.8765 - auc: 0.9441 - mae: 0.1963 - mse: 0.0919 - val_loss: 0.2746 - val_tp: 112.0000 - val_fp: 14.0000 - val_tn: 112.0000 - val_fn: 14.0000 - val_accuracy: 0.8889 - val_precision: 0.8889 - val_auc: 0.9602 - val_mae: 0.1520 - val_mse: 0.0732

Epoch 4/50

46/46 [=====] - 16s 338ms/step - loss: 0.3198 - tp: 1269.0000 - fp: 189.0000 - tn: 1269.0000 - fn: 189.0000 - accuracy: 0.8704 - precision: 0.8704 - auc: 0.9374 - mae: 0.1904 - mse: 0.0979 - val_loss: 0.3726 - val_tp: 107.0000 - val_fp: 19.0000 - val_tn: 107.0000 - val_fn: 19.0000 - val_accuracy: 0.8492 - val_precision: 0.8492 - val_auc: 0.9238 - val_mae: 0.2567 - val_mse: 0.1147

Epoch 5/50

46/46 [=====] - 16s 338ms/step - loss: 0.3149 - tp: 1281.0000 - fp: 177.0000 - tn: 1281.0000 - fn: 177.0000 - accuracy: 0.8786 - precision: 0.8786 - auc: 0.9395 - mae: 0.1881 - mse: 0.0939 - val_loss: 0.2225 - val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 - val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9689 - val_mae: 0.1165 - val_mse: 0.0587

Epoch 6/50

46/46 [=====] - 16s 338ms/step - loss: 0.3507 - tp: 1250.0000 - fp: 208.0000 - tn: 1250.0000 - fn: 208.0000 - accuracy: 0.8573 - precision: 0.8573 - auc: 0.9273 - mae: 0.2080 - mse: 0.1060 - val_loss: 0.2439 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9678 - val_mae: 0.1422 - val_mse: 0.0646

Epoch 7/50

46/46 [=====] - 16s 338ms/step - loss: 0.2579 - tp: 1301.0000 - fp: 157.0000 - tn: 1301.0000 - fn: 157.0000 - accuracy: 0.8923 - precision: 0.8923 - auc: 0.9593 - mae: 0.1553 - mse: 0.0786 - val_loss: 0.2648 - val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 - val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9620 - val_mae: 0.1374 - val_mse: 0.0660

Epoch 8/50

46/46 [=====] - 15s 337ms/step - loss: 0.2473 - tp: 1313.0000 - fp: 145.0000 - tn: 1313.0000 - fn: 145.0000 - accuracy: 0.9005 - precision: 0.9005 - auc: 0.9621 - mae: 0.1512 - mse: 0.0744 - val_loss: 0.2394 - val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 - val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9653 - val_mae: 0.1260 - val_mse: 0.0680

Epoch 9/50

46/46 [=====] - 16s 339ms/step - loss: 0.2763 - tp: 1300.0000 - fp: 158.0000 - tn: 1300.0000 - fn: 158.0000 - accuracy: 0.8916 - precision: 0.8916 - auc: 0.9533 - mae: 0.1702 - mse: 0.0827 - val_loss: 0.4199 - val_tp: 107.0000 - val_fp: 19.0000 - val_tn: 107.0000 - val_fn: 19.0000 - val_accuracy: 0.8492 - val_precision: 0.8492 - val_auc: 0.9273 - val_mae: 0.1730 - val_mse: 0.1141

Epoch 10/50

46/46 [=====] - 16s 340ms/step - loss: 0.2743 - tp: 1295.0000 - fp: 163.0000 - tn: 1295.0000 - fn: 163.0000 - accuracy: 0.8882 - precision: 0.8882 - auc: 0.9553 - mae: 0.1547 - mse: 0.0824 - val_loss: 0.2262 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9686 - val_mae: 0.1446 - val_mse: 0.0661

Epoch 11/50

46/46 [=====] - 16s 339ms/step - loss: 0.2594 - tp: 1303.0000 - fp: 155.0000 - tn: 1303.0000 - fn: 155.0000 - accuracy: 0.8937 - precision: 0.8937 - auc: 0.9584 - mae: 0.1563 - mse: 0.0772 - val_loss: 0.1964 - val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 - val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9747 - val_mae: 0.1137 - val_mse: 0.0562

Epoch 12/50

46/46 [=====] - 16s 337ms/step - loss: 0.2601 - tp: 1306.0000 - fp: 152.0000 - tn: 1306.0000 - fn: 152.0000 - accuracy: 0.8957 - precision: 0.8957 - auc: 0.9587 - mae: 0.1546 - mse: 0.0784 - val_loss: 0.2735 - val_tp: 112.0000 - val_fp: 14.0000 - val_tn: 112.0000 - val_fn: 14.0000 - val_accuracy: 0.8889 - val_precision: 0.8889 - val_auc: 0.9559 - val_mae: 0.1682 - val_mse: 0.0810

Epoch 13/50

46/46 [=====] - 15s 336ms/step - loss: 0.2672 - tp: 1317.0000 - fp: 141.0000 - tn: 1317.0000 - fn: 141.0000 - accuracy: 0.9033 - precision: 0.9033 - auc: 0.9571 - mae: 0.1501 - mse: 0.0772 - val_loss: 0.2232 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9689 - val_mae: 0.1249 - val_mse: 0.0648

Epoch 14/50

46/46 [=====] - 15s 337ms/step - loss: 0.2539 - tp: 1311.0000 - fp: 147.0000 - tn: 1311.0000 - fn: 147.0000 - accuracy: 0.8992 - precision: 0.8992 - auc: 0.9607 - mae: 0.1490 - mse: 0.0753 - val_loss: 0.2395 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9657 - val_mae: 0.1212 - val_mse: 0.0678

Epoch 15/50

46/46 [=====] - 15s 337ms/step - loss: 0.2664 - tp: 1305.0000 - fp: 153.0000 - tn: 1305.0000 - fn: 153.0000 - accuracy: 0.8951 - precision: 0.8951 - auc: 0.9568 - mae: 0.1538 - mse: 0.0786 - val_loss: 0.4520 - val_tp: 102.0000 - val_fp: 24.0000 - val_tn: 102.0000 - val_fn: 24.0000 - val_accuracy: 0.8095 - val_precision: 0.8095 - val_auc: 0.8715 - val_mae: 0.3187 - val_mse: 0.1448

Epoch 16/50

46/46 [=====] - 16s 337ms/step - loss: 0.2807 - tp: 1290.0000 - fp: 168.0000 - tn: 1290.0000 - fn: 168.0000 - accuracy: 0.8848 - precision: 0.8848 - auc: 0.9518 - mae: 0.1732 - mse: 0.0853 - val_loss: 0.2344 - val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 - val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9760 - val_mae: 0.1739 - val_mse: 0.0673

Epoch 17/50

46/46 [=====] - 16s 337ms/step - loss: 0.2591 - tp: 1307.0000 - fp: 151.0000 - tn: 1307.0000 - fn: 151.0000 - accuracy: 0.8964 - precision: 0.8964 - auc: 0.9592 - mae: 0.1517 - mse: 0.0774 - val_loss: 0.4242 - val_tp: 106.0000 - val_fp: 20.0000 - val_tn: 106.0000 - val_fn: 20.0000 - val_accuracy: 0.8413 - val_precision: 0.8413 - val_auc: 0.9222 - val_mae: 0.1901 - val_mse: 0.1115

Epoch 18/50

46/46 [=====] - 16s 337ms/step - loss: 0.2789 - tp: 1298.0000 - fp: 160.0000 - tn: 1298.0000 - fn: 160.0000 - accuracy: 0.8903 - precision: 0.8903 - auc: 0.9532 - mae: 0.1537 - mse: 0.0803 - val_loss: 0.2963 - val_tp: 109.0000 - val_fp: 17.0000 - val_tn: 109.0000 - val_fn: 17.0000 - val_accuracy: 0.8651 - val_precision: 0.8651 - val_auc: 0.9569 - val_mae: 0.1402

```

- val_mse: 0.0846
Epoch 19/50
46/46 [=====] - 16s 338ms/step - loss: 0.2679 - tp:
1304.0000 - fp: 154.0000 - tn: 1304.0000 - fn: 154.0000 - accuracy: 0.8944 -
precision: 0.8944 - auc: 0.9571 - mae: 0.1592 - mse: 0.0784 - val_loss: 0.2944 -
val_tp: 111.0000 - val_fp: 15.0000 - val_tn: 111.0000 - val_fn: 15.0000 -
val_accuracy: 0.8810 - val_precision: 0.8810 - val_auc: 0.9535 - val_mae: 0.1615
- val_mse: 0.0810
Epoch 20/50
46/46 [=====] - 16s 339ms/step - loss: 0.3092 - tp:
1280.0000 - fp: 178.0000 - tn: 1280.0000 - fn: 178.0000 - accuracy: 0.8779 -
precision: 0.8779 - auc: 0.9436 - mae: 0.1757 - mse: 0.0954 - val_loss: 0.2356 -
val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 -
val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9678 - val_mae: 0.1193
- val_mse: 0.0695
Epoch 21/50
46/46 [=====] - 16s 338ms/step - loss: 0.2525 - tp:
1313.0000 - fp: 145.0000 - tn: 1313.0000 - fn: 145.0000 - accuracy: 0.9005 -
precision: 0.9005 - auc: 0.9607 - mae: 0.1498 - mse: 0.0753 - val_loss: 0.2327 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9689 - val_mae: 0.1322
- val_mse: 0.0732
Epoch 22/50
46/46 [=====] - 16s 339ms/step - loss: 0.2509 - tp:
1317.0000 - fp: 141.0000 - tn: 1317.0000 - fn: 141.0000 - accuracy: 0.9033 -
precision: 0.9033 - auc: 0.9617 - mae: 0.1382 - mse: 0.0717 - val_loss: 0.2653 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9598 - val_mae: 0.1175
- val_mse: 0.0672
Epoch 23/50
46/46 [=====] - 16s 342ms/step - loss: 0.2537 - tp:
1314.0000 - fp: 144.0000 - tn: 1314.0000 - fn: 144.0000 - accuracy: 0.9012 -
precision: 0.9012 - auc: 0.9606 - mae: 0.1449 - mse: 0.0753 - val_loss: 0.2878 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9514 - val_mae: 0.1595
- val_mse: 0.0836
Epoch 24/50
46/46 [=====] - 16s 338ms/step - loss: 0.2435 - tp:
1319.0000 - fp: 139.0000 - tn: 1319.0000 - fn: 139.0000 - accuracy: 0.9047 -
precision: 0.9047 - auc: 0.9637 - mae: 0.1378 - mse: 0.0713 - val_loss: 0.2963 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9486 - val_mae: 0.1906
- val_mse: 0.0855
Epoch 25/50
46/46 [=====] - 15s 337ms/step - loss: 0.2874 - tp:
1285.0000 - fp: 173.0000 - tn: 1285.0000 - fn: 173.0000 - accuracy: 0.8813 -
precision: 0.8813 - auc: 0.9512 - mae: 0.1690 - mse: 0.0851 - val_loss: 0.2369 -
val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 -

```

val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9652 - val_mae: 0.1428
- val_mse: 0.0684

Epoch 26/50

46/46 [=====] - 16s 339ms/step - loss: 0.2471 - tp:
1310.0000 - fp: 148.0000 - tn: 1310.0000 - fn: 148.0000 - accuracy: 0.8985 -
precision: 0.8985 - auc: 0.9627 - mae: 0.1412 - mse: 0.0746 - val_loss: 0.2409 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9666 - val_mae: 0.1438
- val_mse: 0.0693

Epoch 27/50

46/46 [=====] - 16s 341ms/step - loss: 0.2507 - tp:
1321.0000 - fp: 137.0000 - tn: 1321.0000 - fn: 137.0000 - accuracy: 0.9060 -
precision: 0.9060 - auc: 0.9623 - mae: 0.1393 - mse: 0.0719 - val_loss: 0.2217 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9703 - val_mae: 0.1069
- val_mse: 0.0605

Epoch 28/50

46/46 [=====] - 16s 341ms/step - loss: 0.2678 - tp:
1297.0000 - fp: 161.0000 - tn: 1297.0000 - fn: 161.0000 - accuracy: 0.8896 -
precision: 0.8896 - auc: 0.9574 - mae: 0.1527 - mse: 0.0799 - val_loss: 0.3336 -
val_tp: 112.0000 - val_fp: 14.0000 - val_tn: 112.0000 - val_fn: 14.0000 -
val_accuracy: 0.8889 - val_precision: 0.8889 - val_auc: 0.9449 - val_mae: 0.1452
- val_mse: 0.0881

Epoch 29/50

46/46 [=====] - 16s 338ms/step - loss: 0.2468 - tp:
1315.0000 - fp: 143.0000 - tn: 1315.0000 - fn: 143.0000 - accuracy: 0.9019 -
precision: 0.9019 - auc: 0.9625 - mae: 0.1447 - mse: 0.0723 - val_loss: 0.1973 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9786 - val_mae: 0.1327
- val_mse: 0.0591

Epoch 30/50

46/46 [=====] - 16s 338ms/step - loss: 0.2197 - tp:
1335.0000 - fp: 123.0000 - tn: 1335.0000 - fn: 123.0000 - accuracy: 0.9156 -
precision: 0.9156 - auc: 0.9711 - mae: 0.1309 - mse: 0.0644 - val_loss: 0.2448 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9667 - val_mae: 0.1521
- val_mse: 0.0695

Epoch 31/50

46/46 [=====] - 16s 337ms/step - loss: 0.2140 - tp:
1331.0000 - fp: 127.0000 - tn: 1331.0000 - fn: 127.0000 - accuracy: 0.9129 -
precision: 0.9129 - auc: 0.9725 - mae: 0.1255 - mse: 0.0630 - val_loss: 0.2556 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9660 - val_mae: 0.1794
- val_mse: 0.0777

Epoch 32/50

46/46 [=====] - 16s 338ms/step - loss: 0.2312 - tp:
1321.0000 - fp: 137.0000 - tn: 1321.0000 - fn: 137.0000 - accuracy: 0.9060 -
precision: 0.9060 - auc: 0.9676 - mae: 0.1367 - mse: 0.0695 - val_loss: 0.2516 -

val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 -
val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9616 - val_mae: 0.1261
- val_mse: 0.0647

Epoch 33/50

46/46 [=====] - 16s 337ms/step - loss: 0.2324 - tp:
1326.0000 - fp: 132.0000 - tn: 1326.0000 - fn: 132.0000 - accuracy: 0.9095 -
precision: 0.9095 - auc: 0.9667 - mae: 0.1327 - mse: 0.0689 - val_loss: 0.1911 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9782 - val_mae: 0.1135
- val_mse: 0.0587

Epoch 34/50

46/46 [=====] - 16s 339ms/step - loss: 0.2318 - tp:
1332.0000 - fp: 126.0000 - tn: 1332.0000 - fn: 126.0000 - accuracy: 0.9136 -
precision: 0.9136 - auc: 0.9683 - mae: 0.1313 - mse: 0.0673 - val_loss: 0.2536 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9640 - val_mae: 0.1669
- val_mse: 0.0729

Epoch 35/50

46/46 [=====] - 16s 338ms/step - loss: 0.4999 - tp:
1129.0000 - fp: 329.0000 - tn: 1129.0000 - fn: 329.0000 - accuracy: 0.7743 -
precision: 0.7743 - auc: 0.8484 - mae: 0.3105 - mse: 0.1599 - val_loss: 0.5549 -
val_tp: 108.0000 - val_fp: 18.0000 - val_tn: 108.0000 - val_fn: 18.0000 -
val_accuracy: 0.8571 - val_precision: 0.8571 - val_auc: 0.8745 - val_mae: 0.2766
- val_mse: 0.1474

Epoch 36/50

46/46 [=====] - 15s 337ms/step - loss: 0.4832 - tp:
1153.0000 - fp: 305.0000 - tn: 1153.0000 - fn: 305.0000 - accuracy: 0.7908 -
precision: 0.7908 - auc: 0.8562 - mae: 0.3000 - mse: 0.1547 - val_loss: 0.4154 -
val_tp: 108.0000 - val_fp: 18.0000 - val_tn: 108.0000 - val_fn: 18.0000 -
val_accuracy: 0.8571 - val_precision: 0.8571 - val_auc: 0.9145 - val_mae: 0.2922
- val_mse: 0.1303

Epoch 37/50

46/46 [=====] - 16s 337ms/step - loss: 0.4453 - tp:
1184.0000 - fp: 274.0000 - tn: 1184.0000 - fn: 274.0000 - accuracy: 0.8121 -
precision: 0.8121 - auc: 0.8795 - mae: 0.2784 - mse: 0.1409 - val_loss: 0.4285 -
val_tp: 104.0000 - val_fp: 22.0000 - val_tn: 104.0000 - val_fn: 22.0000 -
val_accuracy: 0.8254 - val_precision: 0.8254 - val_auc: 0.9086 - val_mae: 0.2352
- val_mse: 0.1218

Epoch 38/50

46/46 [=====] - 16s 337ms/step - loss: 0.4062 - tp:
1249.0000 - fp: 209.0000 - tn: 1249.0000 - fn: 209.0000 - accuracy: 0.8567 -
precision: 0.8567 - auc: 0.9070 - mae: 0.2235 - mse: 0.1184 - val_loss: 0.3565 -
val_tp: 106.0000 - val_fp: 20.0000 - val_tn: 106.0000 - val_fn: 20.0000 -
val_accuracy: 0.8413 - val_precision: 0.8413 - val_auc: 0.9238 - val_mae: 0.2337
- val_mse: 0.1105

Epoch 39/50

46/46 [=====] - 15s 336ms/step - loss: 0.3689 - tp:
1246.0000 - fp: 212.0000 - tn: 1246.0000 - fn: 212.0000 - accuracy: 0.8546 -

precision: 0.8546 - auc: 0.9182 - mae: 0.2110 - mse: 0.1105 - val_loss: 0.2857 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9523 - val_mae: 0.1426
- val_mse: 0.0733

Epoch 40/50

46/46 [=====] - 15s 337ms/step - loss: 0.3129 - tp:
1286.0000 - fp: 172.0000 - tn: 1286.0000 - fn: 172.0000 - accuracy: 0.8820 -
precision: 0.8820 - auc: 0.9468 - mae: 0.1580 - mse: 0.0859 - val_loss: 0.3209 -
val_tp: 107.0000 - val_fp: 19.0000 - val_tn: 107.0000 - val_fn: 19.0000 -
val_accuracy: 0.8492 - val_precision: 0.8492 - val_auc: 0.9405 - val_mae: 0.1820
- val_mse: 0.1016

Epoch 41/50

46/46 [=====] - 16s 337ms/step - loss: 0.3089 - tp:
1290.0000 - fp: 168.0000 - tn: 1290.0000 - fn: 168.0000 - accuracy: 0.8848 -
precision: 0.8848 - auc: 0.9477 - mae: 0.1538 - mse: 0.0894 - val_loss: 0.2159 -
val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 -
val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9710 - val_mae: 0.1162
- val_mse: 0.0599

Epoch 42/50

46/46 [=====] - 15s 337ms/step - loss: 0.2365 - tp:
1322.0000 - fp: 136.0000 - tn: 1322.0000 - fn: 136.0000 - accuracy: 0.9067 -
precision: 0.9067 - auc: 0.9662 - mae: 0.1312 - mse: 0.0701 - val_loss: 0.2409 -
val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 -
val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9646 - val_mae: 0.1387
- val_mse: 0.0738

Epoch 43/50

46/46 [=====] - 15s 337ms/step - loss: 0.2628 - tp:
1304.0000 - fp: 154.0000 - tn: 1304.0000 - fn: 154.0000 - accuracy: 0.8944 -
precision: 0.8944 - auc: 0.9589 - mae: 0.1458 - mse: 0.0779 - val_loss: 0.2772 -
val_tp: 110.0000 - val_fp: 16.0000 - val_tn: 110.0000 - val_fn: 16.0000 -
val_accuracy: 0.8730 - val_precision: 0.8730 - val_auc: 0.9544 - val_mae: 0.1520
- val_mse: 0.0836

Epoch 44/50

46/46 [=====] - 15s 337ms/step - loss: 0.2094 - tp:
1339.0000 - fp: 119.0000 - tn: 1339.0000 - fn: 119.0000 - accuracy: 0.9184 -
precision: 0.9184 - auc: 0.9726 - mae: 0.1180 - mse: 0.0591 - val_loss: 0.2771 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9528 - val_mae: 0.1416
- val_mse: 0.0790

Epoch 45/50

46/46 [=====] - 16s 338ms/step - loss: 0.2531 - tp:
1316.0000 - fp: 142.0000 - tn: 1316.0000 - fn: 142.0000 - accuracy: 0.9026 -
precision: 0.9026 - auc: 0.9624 - mae: 0.1370 - mse: 0.0725 - val_loss: 0.2379 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9679 - val_mae: 0.1522
- val_mse: 0.0650

Epoch 46/50

46/46 [=====] - 16s 341ms/step - loss: 0.2234 - tp:

1330.0000 - fp: 128.0000 - tn: 1330.0000 - fn: 128.0000 - accuracy: 0.9122 -
precision: 0.9122 - auc: 0.9698 - mae: 0.1253 - mse: 0.0652 - val_loss: 0.2621 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9703 - val_mae: 0.0978
- val_mse: 0.0608

Epoch 47/50

46/46 [=====] - 16s 338ms/step - loss: 0.2217 - tp:
1333.0000 - fp: 125.0000 - tn: 1333.0000 - fn: 125.0000 - accuracy: 0.9143 -
precision: 0.9143 - auc: 0.9712 - mae: 0.1208 - mse: 0.0637 - val_loss: 0.2922 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9516 - val_mae: 0.1557
- val_mse: 0.0714

Epoch 48/50

46/46 [=====] - 16s 337ms/step - loss: 0.2050 - tp:
1332.0000 - fp: 126.0000 - tn: 1332.0000 - fn: 126.0000 - accuracy: 0.9136 -
precision: 0.9136 - auc: 0.9741 - mae: 0.1192 - mse: 0.0612 - val_loss: 0.1921 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9781 - val_mae: 0.1091
- val_mse: 0.0571

Epoch 49/50

46/46 [=====] - 15s 336ms/step - loss: 0.2267 - tp:
1325.0000 - fp: 133.0000 - tn: 1325.0000 - fn: 133.0000 - accuracy: 0.9088 -
precision: 0.9088 - auc: 0.9697 - mae: 0.1235 - mse: 0.0667 - val_loss: 0.2111 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9727 - val_mae: 0.1363
- val_mse: 0.0596

Epoch 50/50

46/46 [=====] - 15s 337ms/step - loss: 0.2178 - tp:
1334.0000 - fp: 124.0000 - tn: 1334.0000 - fn: 124.0000 - accuracy: 0.9150 -
precision: 0.9150 - auc: 0.9713 - mae: 0.1209 - mse: 0.0620 - val_loss: 0.2280 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9657 - val_mae: 0.1104
- val_mse: 0.0590

1.1.4 6. Write history and plot graphs 1

```
[11]: def print_history(history):  
    acc = history.history['accuracy']  
    val_acc = history.history['val_accuracy']  
  
    loss = history.history['loss']  
    val_loss = history.history['val_loss']  
  
    auc = history.history['auc']  
    val_auc = history.history['val_auc']  
  
    plt.figure(figsize=(8, 8))
```

```


plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),1])
plt.title('Training and Validation Accuracy')

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.ylabel('Cross Entropy')
plt.ylim([0,1.0])
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(auc, label='Training Auc')
plt.plot(val_auc, label='Validation Auc')
plt.legend(loc='upper right')
plt.ylabel('Auc')
plt.ylim([0,2.0])
plt.title('Training and Validation Auc')
plt.xlabel('epoch')
plt.show()

print_history(history1)

```





1.1.5 7. Evaluate the model 1

```
[12]: def print_model_evaluation(model):
    result = model.evaluate(test_dataset)
    metrics = ["loss", "tp", "fp", "tn", "fn", "accuracy", "precision", "auc",
               ↪ "mae", "mse"]
```



```

    for i in range(len(result)):
        print("{} : {}".format(metrics[i], round(result[i], 3)))

print_model_evaluation(model1);

```

```

714/714 [=====] - 3s 5ms/step - loss: 0.3188 - tp:
624.0000 - fp: 90.0000 - tn: 624.0000 - fn: 90.0000 - accuracy: 0.8739 -
precision: 0.8739 - auc: 0.9494 - mae: 0.1503 - mse: 0.0966
loss : 0.319
tp : 624.0
fp : 90.0
tn : 624.0
fn : 90.0
accuracy : 0.874
precision : 0.874
auc : 0.949
mae : 0.15
mse : 0.097

```

1.1.6 8. Predict with model 1

```

[33]: def predict_and_print_roc(model):
    #Retrieve one batch of images from the test set
    train_dataset.reset()
    image_batch, label_batch = train_dataset.next()

    # Predict the images from the batch
    predictions = model.predict(image_batch)

    # Visualise the actual value and the prediction in numerical form
    probs = list(np.argmax(x) for x in predictions)
    labels = list(np.argmax(x) for x in label_batch)
    print('Label Predictions:\n', probs)
    print('Real Labels:\n', labels)

    confusion = confusion_matrix(labels, probs)
    print('\nConfusion Matrix:')
    print(confusion)

    print('\nAccuracy: {:.2f}\n'.format(accuracy_score(labels, probs)))

    #Print the first 9 Images from the batch and the estimated prediction.
    plt.figure(figsize=(10, 10))
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(image_batch[i])
        plt.title(class_names[probs[i]])

```

```

plt.axis("off")

#Calculate the roc curve
fpr, tpr, _ = skmetrics.roc_curve(label_batch.ravel(), predictions.ravel())
roc_auc = skmetrics.auc(fpr, tpr)

#Display ROC curve and the AUC
plt.figure(figsize=(5, 5))
plt.title('ROC')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

predict_and_print_roc(model1);

```

Label Predictions:

[0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1,
1, 0, 0, 1, 1, 0]

Real Labels:

[0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1,
1, 0, 0, 1, 1, 0]

Confusion Matrix:

```
[[19  0]
 [ 1 12]]
```

Accuracy: 0.97

Apple__healthy



Apple__healthy



Apple__unhealthy



Apple__healthy



Apple__healthy



Apple__healthy



Apple__healthy




Apple__unhealthy



Apple__unhealthy





1.1.7 3. Create and compile model 2

This model uses compared to the previous model three convolution layers and two dense layers. The value of the neurons are decreased at each layer.

```
[14]: model2 = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(128, (3, 3), input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(32, (3, 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(8, activation='relu', use_bias=False),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```

```
[15]: opt = tf.keras.optimizers.Adam(lr=LEARNING_RATE)
model2.compile(optimizer=opt,
```

```
loss='categorical_crossentropy',  
metrics=METRICS)
```

1.1.8 4. Display model structure 2

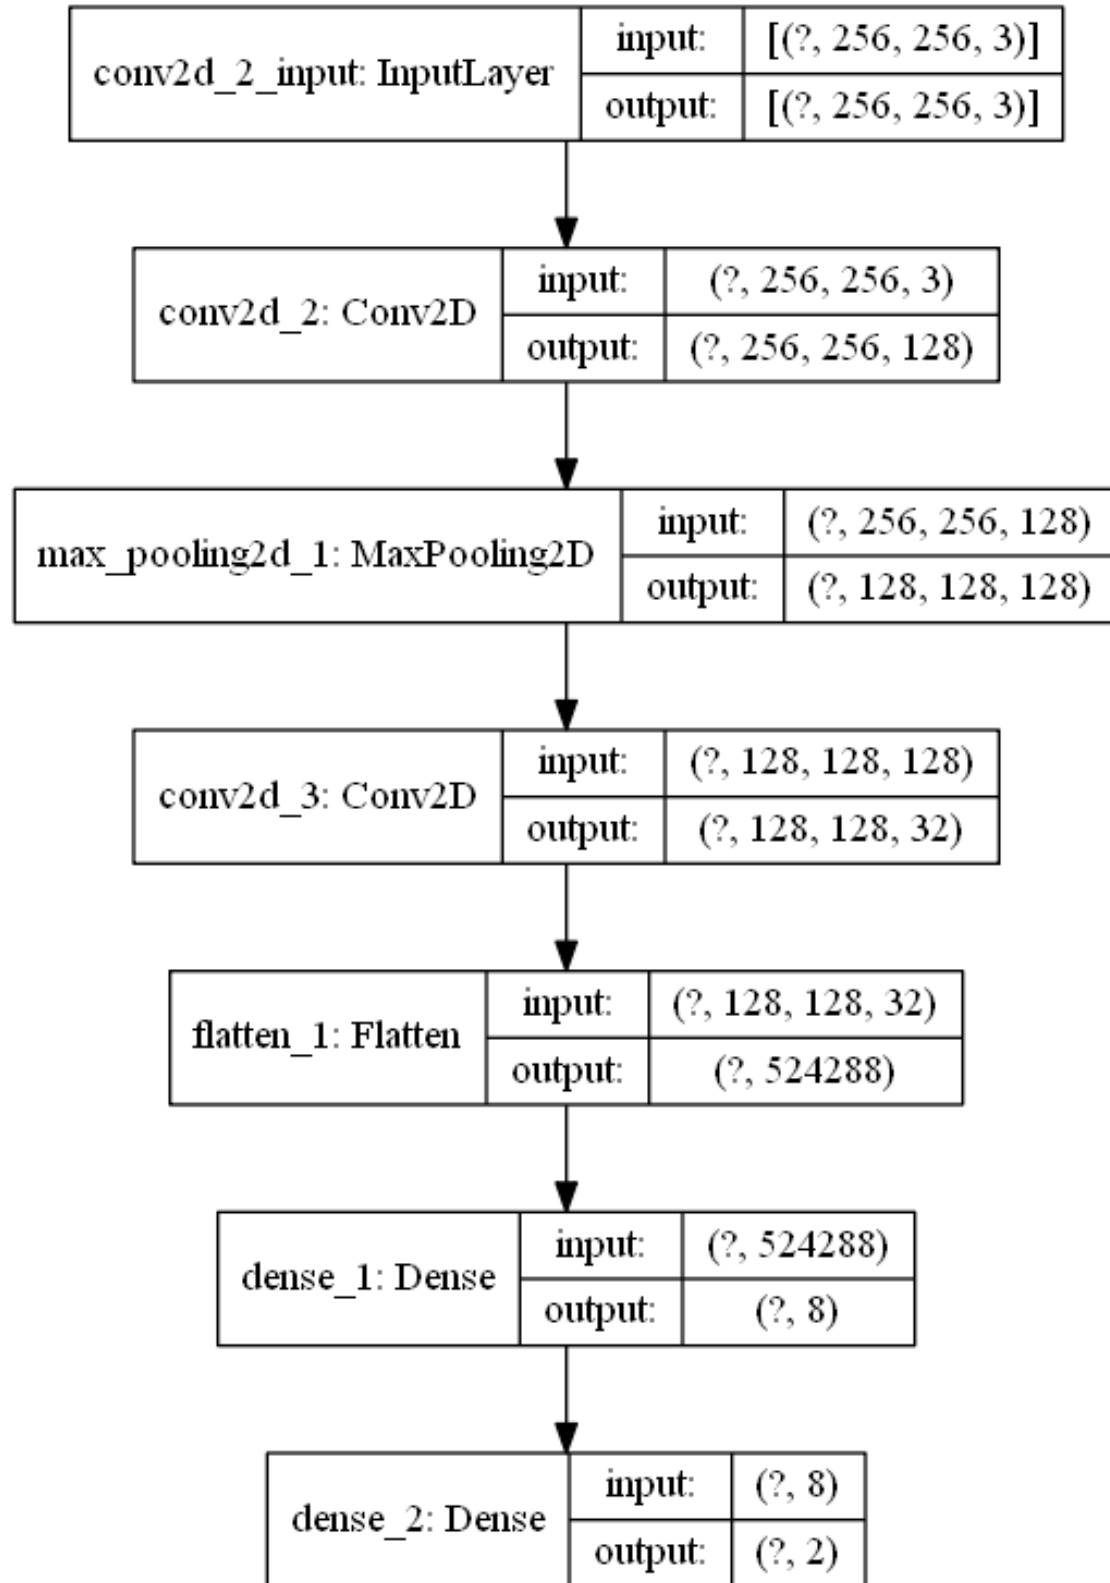
```
[16]: model2.summary()  
      tf.keras.utils.plot_model(model2, show_shapes=True)
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|--------------------------------|-----------------------|---------|
| conv2d_2 (Conv2D) | (None, 256, 256, 128) | 3456 |
| max_pooling2d_1 (MaxPooling2D) | (None, 128, 128, 128) | 0 |
| conv2d_3 (Conv2D) | (None, 128, 128, 32) | 36864 |
| flatten_1 (Flatten) | (None, 524288) | 0 |
| dense_1 (Dense) | (None, 8) | 4194304 |
| dense_2 (Dense) | (None, 2) | 18 |

Total params: 4,234,642
Trainable params: 4,234,642
Non-trainable params: 0

```
[16]:
```



1.1.9 5. Training model 2

```
[17]: history2 = train_model(model2, "Model_2")
```

Epoch 1/50

```
46/46 [=====] - 16s 357ms/step - loss: 0.8398 - tp:
1539.0000 - fp: 632.0000 - tn: 1540.0000 - fn: 633.0000 - accuracy: 0.7086 -
precision: 0.7089 - auc: 0.7869 - mae: 0.3603 - mse: 0.1921 - val_loss: 0.5052 -
val_tp: 103.0000 - val_fp: 23.0000 - val_tn: 103.0000 - val_fn: 23.0000 -
val_accuracy: 0.8175 - val_precision: 0.8175 - val_auc: 0.8818 - val_mae: 0.3769
- val_mse: 0.1639
```

Epoch 2/50

```
46/46 [=====] - 16s 337ms/step - loss: 0.4783 - tp:
1178.0000 - fp: 280.0000 - tn: 1178.0000 - fn: 280.0000 - accuracy: 0.8080 -
precision: 0.8080 - auc: 0.8622 - mae: 0.3090 - mse: 0.1499 - val_loss: 0.3890 -
val_tp: 106.0000 - val_fp: 20.0000 - val_tn: 106.0000 - val_fn: 20.0000 -
val_accuracy: 0.8413 - val_precision: 0.8413 - val_auc: 0.9121 - val_mae: 0.2112
- val_mse: 0.1145
```

Epoch 3/50

```
46/46 [=====] - 15s 335ms/step - loss: 0.3931 - tp:
1223.0000 - fp: 235.0000 - tn: 1223.0000 - fn: 235.0000 - accuracy: 0.8388 -
precision: 0.8388 - auc: 0.9075 - mae: 0.2422 - mse: 0.1192 - val_loss: 0.2747 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9609 - val_mae: 0.2014
- val_mse: 0.0779
```

Epoch 4/50

```
46/46 [=====] - 16s 338ms/step - loss: 0.3778 - tp:
1222.0000 - fp: 236.0000 - tn: 1222.0000 - fn: 236.0000 - accuracy: 0.8381 -
precision: 0.8381 - auc: 0.9133 - mae: 0.2372 - mse: 0.1164 - val_loss: 0.3615 -
val_tp: 106.0000 - val_fp: 20.0000 - val_tn: 106.0000 - val_fn: 20.0000 -
val_accuracy: 0.8413 - val_precision: 0.8413 - val_auc: 0.9323 - val_mae: 0.2656
- val_mse: 0.1112
```

Epoch 5/50

```
46/46 [=====] - 15s 336ms/step - loss: 0.3617 - tp:
1237.0000 - fp: 221.0000 - tn: 1237.0000 - fn: 221.0000 - accuracy: 0.8484 -
precision: 0.8484 - auc: 0.9203 - mae: 0.2201 - mse: 0.1100 - val_loss: 0.4808 -
val_tp: 94.0000 - val_fp: 32.0000 - val_tn: 94.0000 - val_fn: 32.0000 -
val_accuracy: 0.7460 - val_precision: 0.7460 - val_auc: 0.8553 - val_mae: 0.2768
- val_mse: 0.1641
```

Epoch 6/50

```
46/46 [=====] - 15s 336ms/step - loss: 0.4156 - tp:
1189.0000 - fp: 269.0000 - tn: 1189.0000 - fn: 269.0000 - accuracy: 0.8155 -
precision: 0.8155 - auc: 0.8952 - mae: 0.2515 - mse: 0.1299 - val_loss: 0.2465 -
val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 -
val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9707 - val_mae: 0.1823
- val_mse: 0.0692
```

Epoch 7/50

```
46/46 [=====] - 16s 337ms/step - loss: 0.4017 - tp:
```

1217.0000 - fp: 241.0000 - tn: 1217.0000 - fn: 241.0000 - accuracy: 0.8347 -
precision: 0.8347 - auc: 0.9033 - mae: 0.2448 - mse: 0.1230 - val_loss: 0.2629 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9718 - val_mae: 0.2012
- val_mse: 0.0764

Epoch 8/50

46/46 [=====] - 16s 337ms/step - loss: 0.3391 - tp:
1259.0000 - fp: 199.0000 - tn: 1259.0000 - fn: 199.0000 - accuracy: 0.8635 -
precision: 0.8635 - auc: 0.9302 - mae: 0.2109 - mse: 0.1026 - val_loss: 0.2602 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9592 - val_mae: 0.1683
- val_mse: 0.0771

Epoch 9/50

46/46 [=====] - 15s 336ms/step - loss: 0.3346 - tp:
1272.0000 - fp: 186.0000 - tn: 1272.0000 - fn: 186.0000 - accuracy: 0.8724 -
precision: 0.8724 - auc: 0.9326 - mae: 0.2013 - mse: 0.0996 - val_loss: 0.2840 -
val_tp: 108.0000 - val_fp: 18.0000 - val_tn: 108.0000 - val_fn: 18.0000 -
val_accuracy: 0.8571 - val_precision: 0.8571 - val_auc: 0.9499 - val_mae: 0.1718
- val_mse: 0.0935

Epoch 10/50

46/46 [=====] - 15s 335ms/step - loss: 0.3758 - tp:
1228.0000 - fp: 230.0000 - tn: 1228.0000 - fn: 230.0000 - accuracy: 0.8422 -
precision: 0.8422 - auc: 0.9163 - mae: 0.2221 - mse: 0.1136 - val_loss: 0.3727 -
val_tp: 110.0000 - val_fp: 16.0000 - val_tn: 110.0000 - val_fn: 16.0000 -
val_accuracy: 0.8730 - val_precision: 0.8730 - val_auc: 0.9259 - val_mae: 0.1660
- val_mse: 0.1002

Epoch 11/50

46/46 [=====] - 16s 337ms/step - loss: 0.3782 - tp:
1228.0000 - fp: 230.0000 - tn: 1228.0000 - fn: 230.0000 - accuracy: 0.8422 -
precision: 0.8422 - auc: 0.9147 - mae: 0.2322 - mse: 0.1147 - val_loss: 0.2341 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9678 - val_mae: 0.1438
- val_mse: 0.0672

Epoch 12/50

46/46 [=====] - 15s 336ms/step - loss: 0.3143 - tp:
1275.0000 - fp: 183.0000 - tn: 1275.0000 - fn: 183.0000 - accuracy: 0.8745 -
precision: 0.8745 - auc: 0.9403 - mae: 0.1911 - mse: 0.0936 - val_loss: 0.2696 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9560 - val_mae: 0.1497
- val_mse: 0.0801

Epoch 13/50

46/46 [=====] - 15s 336ms/step - loss: 0.3255 - tp:
1270.0000 - fp: 188.0000 - tn: 1270.0000 - fn: 188.0000 - accuracy: 0.8711 -
precision: 0.8711 - auc: 0.9358 - mae: 0.2004 - mse: 0.0980 - val_loss: 0.2789 -
val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 -
val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9535 - val_mae: 0.1661
- val_mse: 0.0818

Epoch 14/50

46/46 [=====] - 16s 338ms/step - loss: 0.3405 - tp: 1252.0000 - fp: 206.0000 - tn: 1252.0000 - fn: 206.0000 - accuracy: 0.8587 - precision: 0.8587 - auc: 0.9308 - mae: 0.2023 - mse: 0.1029 - val_loss: 0.2067 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9737 - val_mae: 0.1297 - val_mse: 0.0583

Epoch 15/50

46/46 [=====] - 15s 336ms/step - loss: 0.3265 - tp: 1271.0000 - fp: 187.0000 - tn: 1271.0000 - fn: 187.0000 - accuracy: 0.8717 - precision: 0.8717 - auc: 0.9359 - mae: 0.2002 - mse: 0.0974 - val_loss: 0.1776 - val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 - val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9836 - val_mae: 0.1248 - val_mse: 0.0526

Epoch 16/50

46/46 [=====] - 15s 335ms/step - loss: 0.3243 - tp: 1263.0000 - fp: 195.0000 - tn: 1263.0000 - fn: 195.0000 - accuracy: 0.8663 - precision: 0.8663 - auc: 0.9364 - mae: 0.1963 - mse: 0.0985 - val_loss: 0.2259 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9691 - val_mae: 0.1452 - val_mse: 0.0679

Epoch 17/50

46/46 [=====] - 15s 337ms/step - loss: 0.3145 - tp: 1289.0000 - fp: 169.0000 - tn: 1289.0000 - fn: 169.0000 - accuracy: 0.8841 - precision: 0.8841 - auc: 0.9401 - mae: 0.1920 - mse: 0.0937 - val_loss: 0.2119 - val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 - val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9695 - val_mae: 0.1265 - val_mse: 0.0596

Epoch 18/50

46/46 [=====] - 15s 336ms/step - loss: 0.3350 - tp: 1255.0000 - fp: 203.0000 - tn: 1255.0000 - fn: 203.0000 - accuracy: 0.8608 - precision: 0.8608 - auc: 0.9335 - mae: 0.1974 - mse: 0.1006 - val_loss: 0.2512 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9692 - val_mae: 0.1822 - val_mse: 0.0725

Epoch 19/50

46/46 [=====] - 15s 336ms/step - loss: 0.3100 - tp: 1288.0000 - fp: 170.0000 - tn: 1288.0000 - fn: 170.0000 - accuracy: 0.8834 - precision: 0.8834 - auc: 0.9423 - mae: 0.1869 - mse: 0.0910 - val_loss: 0.6489 - val_tp: 86.0000 - val_fp: 40.0000 - val_tn: 86.0000 - val_fn: 40.0000 - val_accuracy: 0.6825 - val_precision: 0.6825 - val_auc: 0.8376 - val_mae: 0.2780 - val_mse: 0.1991

Epoch 20/50

46/46 [=====] - 15s 337ms/step - loss: 0.3894 - tp: 1216.0000 - fp: 242.0000 - tn: 1216.0000 - fn: 242.0000 - accuracy: 0.8340 - precision: 0.8340 - auc: 0.9104 - mae: 0.2232 - mse: 0.1181 - val_loss: 0.4272 - val_tp: 100.0000 - val_fp: 26.0000 - val_tn: 100.0000 - val_fn: 26.0000 - val_accuracy: 0.7937 - val_precision: 0.7937 - val_auc: 0.9030 - val_mae: 0.3156 - val_mse: 0.1363

Epoch 21/50

46/46 [=====] - 15s 337ms/step - loss: 0.4102 - tp: 1190.0000 - fp: 268.0000 - tn: 1190.0000 - fn: 268.0000 - accuracy: 0.8162 - precision: 0.8162 - auc: 0.8994 - mae: 0.2470 - mse: 0.1264 - val_loss: 0.2247 - val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 - val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9750 - val_mae: 0.1631 - val_mse: 0.0645

Epoch 22/50

46/46 [=====] - 15s 337ms/step - loss: 0.4122 - tp: 1178.0000 - fp: 280.0000 - tn: 1178.0000 - fn: 280.0000 - accuracy: 0.8080 - precision: 0.8080 - auc: 0.8956 - mae: 0.2705 - mse: 0.1334 - val_loss: 0.2553 - val_tp: 112.0000 - val_fp: 14.0000 - val_tn: 112.0000 - val_fn: 14.0000 - val_accuracy: 0.8889 - val_precision: 0.8889 - val_auc: 0.9595 - val_mae: 0.1517 - val_mse: 0.0764

Epoch 23/50

46/46 [=====] - 15s 336ms/step - loss: 0.3001 - tp: 1299.0000 - fp: 159.0000 - tn: 1299.0000 - fn: 159.0000 - accuracy: 0.8909 - precision: 0.8909 - auc: 0.9466 - mae: 0.1781 - mse: 0.0864 - val_loss: 0.2006 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9769 - val_mae: 0.1328 - val_mse: 0.0607

Epoch 24/50

46/46 [=====] - 15s 336ms/step - loss: 0.3766 - tp: 1217.0000 - fp: 241.0000 - tn: 1217.0000 - fn: 241.0000 - accuracy: 0.8347 - precision: 0.8347 - auc: 0.9158 - mae: 0.2188 - mse: 0.1150 - val_loss: 0.2136 - val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 - val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9744 - val_mae: 0.1435 - val_mse: 0.0637

Epoch 25/50

46/46 [=====] - 15s 336ms/step - loss: 0.2960 - tp: 1299.0000 - fp: 159.0000 - tn: 1299.0000 - fn: 159.0000 - accuracy: 0.8909 - precision: 0.8909 - auc: 0.9466 - mae: 0.1813 - mse: 0.0874 - val_loss: 0.2084 - val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 - val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9728 - val_mae: 0.1315 - val_mse: 0.0601

Epoch 26/50

46/46 [=====] - 15s 337ms/step - loss: 0.2922 - tp: 1292.0000 - fp: 166.0000 - tn: 1292.0000 - fn: 166.0000 - accuracy: 0.8861 - precision: 0.8861 - auc: 0.9484 - mae: 0.1769 - mse: 0.0864 - val_loss: 0.2885 - val_tp: 109.0000 - val_fp: 17.0000 - val_tn: 109.0000 - val_fn: 17.0000 - val_accuracy: 0.8651 - val_precision: 0.8651 - val_auc: 0.9517 - val_mae: 0.1998 - val_mse: 0.0906

Epoch 27/50

46/46 [=====] - 15s 336ms/step - loss: 0.3076 - tp: 1281.0000 - fp: 177.0000 - tn: 1281.0000 - fn: 177.0000 - accuracy: 0.8786 - precision: 0.8786 - auc: 0.9422 - mae: 0.1858 - mse: 0.0919 - val_loss: 0.2063 - val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 - val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9747 - val_mae: 0.1389

```

- val_mse: 0.0587
Epoch 28/50
46/46 [=====] - 15s 336ms/step - loss: 0.2852 - tp:
1279.0000 - fp: 179.0000 - tn: 1279.0000 - fn: 179.0000 - accuracy: 0.8772 -
precision: 0.8772 - auc: 0.9502 - mae: 0.1775 - mse: 0.0885 - val_loss: 0.1913 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9789 - val_mae: 0.1259
- val_mse: 0.0545
Epoch 29/50
46/46 [=====] - 15s 336ms/step - loss: 0.2818 - tp:
1299.0000 - fp: 159.0000 - tn: 1299.0000 - fn: 159.0000 - accuracy: 0.8909 -
precision: 0.8909 - auc: 0.9517 - mae: 0.1693 - mse: 0.0842 - val_loss: 0.2666 -
val_tp: 112.0000 - val_fp: 14.0000 - val_tn: 112.0000 - val_fn: 14.0000 -
val_accuracy: 0.8889 - val_precision: 0.8889 - val_auc: 0.9569 - val_mae: 0.1515
- val_mse: 0.0836
Epoch 30/50
46/46 [=====] - 15s 336ms/step - loss: 0.2992 - tp:
1281.0000 - fp: 177.0000 - tn: 1281.0000 - fn: 177.0000 - accuracy: 0.8786 -
precision: 0.8786 - auc: 0.9461 - mae: 0.1743 - mse: 0.0903 - val_loss: 0.2772 -
val_tp: 108.0000 - val_fp: 18.0000 - val_tn: 108.0000 - val_fn: 18.0000 -
val_accuracy: 0.8571 - val_precision: 0.8571 - val_auc: 0.9538 - val_mae: 0.2005
- val_mse: 0.0898
Epoch 31/50
46/46 [=====] - 15s 336ms/step - loss: 0.3136 - tp:
1272.0000 - fp: 186.0000 - tn: 1272.0000 - fn: 186.0000 - accuracy: 0.8724 -
precision: 0.8724 - auc: 0.9401 - mae: 0.1950 - mse: 0.0963 - val_loss: 0.2658 -
val_tp: 110.0000 - val_fp: 16.0000 - val_tn: 110.0000 - val_fn: 16.0000 -
val_accuracy: 0.8730 - val_precision: 0.8730 - val_auc: 0.9562 - val_mae: 0.1762
- val_mse: 0.0850
Epoch 32/50
46/46 [=====] - 15s 336ms/step - loss: 0.3012 - tp:
1286.0000 - fp: 172.0000 - tn: 1286.0000 - fn: 172.0000 - accuracy: 0.8820 -
precision: 0.8820 - auc: 0.9457 - mae: 0.1767 - mse: 0.0902 - val_loss: 0.1955 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9772 - val_mae: 0.1308
- val_mse: 0.0552
Epoch 33/50
46/46 [=====] - 15s 337ms/step - loss: 0.2937 - tp:
1283.0000 - fp: 175.0000 - tn: 1283.0000 - fn: 175.0000 - accuracy: 0.8800 -
precision: 0.8800 - auc: 0.9484 - mae: 0.1844 - mse: 0.0878 - val_loss: 0.1881 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9793 - val_mae: 0.1270
- val_mse: 0.0577
Epoch 34/50
46/46 [=====] - 16s 340ms/step - loss: 0.2705 - tp:
1293.0000 - fp: 165.0000 - tn: 1293.0000 - fn: 165.0000 - accuracy: 0.8868 -
precision: 0.8868 - auc: 0.9560 - mae: 0.1653 - mse: 0.0805 - val_loss: 0.2080 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -

```

val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9783 - val_mae: 0.1524
- val_mse: 0.0627

Epoch 35/50

46/46 [=====] - 15s 335ms/step - loss: 0.2640 - tp:
1300.0000 - fp: 158.0000 - tn: 1300.0000 - fn: 158.0000 - accuracy: 0.8916 -
precision: 0.8916 - auc: 0.9574 - mae: 0.1612 - mse: 0.0792 - val_loss: 0.2017 -
val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 -
val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9785 - val_mae: 0.1425
- val_mse: 0.0611

Epoch 36/50

46/46 [=====] - 15s 336ms/step - loss: 0.2697 - tp:
1305.0000 - fp: 153.0000 - tn: 1305.0000 - fn: 153.0000 - accuracy: 0.8951 -
precision: 0.8951 - auc: 0.9559 - mae: 0.1649 - mse: 0.0813 - val_loss: 0.1875 -
val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 -
val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9791 - val_mae: 0.1182
- val_mse: 0.0567

Epoch 37/50

46/46 [=====] - 15s 335ms/step - loss: 0.2768 - tp:
1293.0000 - fp: 165.0000 - tn: 1293.0000 - fn: 165.0000 - accuracy: 0.8868 -
precision: 0.8868 - auc: 0.9536 - mae: 0.1617 - mse: 0.0832 - val_loss: 0.2027 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9754 - val_mae: 0.1188
- val_mse: 0.0604

Epoch 38/50

46/46 [=====] - 15s 336ms/step - loss: 0.2664 - tp:
1298.0000 - fp: 160.0000 - tn: 1298.0000 - fn: 160.0000 - accuracy: 0.8903 -
precision: 0.8903 - auc: 0.9575 - mae: 0.1671 - mse: 0.0801 - val_loss: 0.2131 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9715 - val_mae: 0.1263
- val_mse: 0.0640

Epoch 39/50

46/46 [=====] - 15s 337ms/step - loss: 0.2792 - tp:
1302.0000 - fp: 156.0000 - tn: 1302.0000 - fn: 156.0000 - accuracy: 0.8930 -
precision: 0.8930 - auc: 0.9527 - mae: 0.1625 - mse: 0.0818 - val_loss: 0.2298 -
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9682 - val_mae: 0.1411
- val_mse: 0.0713

Epoch 40/50

46/46 [=====] - 15s 336ms/step - loss: 0.2619 - tp:
1309.0000 - fp: 149.0000 - tn: 1309.0000 - fn: 149.0000 - accuracy: 0.8978 -
precision: 0.8978 - auc: 0.9577 - mae: 0.1540 - mse: 0.0766 - val_loss: 0.1848 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9807 - val_mae: 0.1281
- val_mse: 0.0551

Epoch 41/50

46/46 [=====] - 15s 337ms/step - loss: 0.2389 - tp:
1319.0000 - fp: 139.0000 - tn: 1319.0000 - fn: 139.0000 - accuracy: 0.9047 -
precision: 0.9047 - auc: 0.9650 - mae: 0.1495 - mse: 0.0715 - val_loss: 0.1717 -

val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9808 - val_mae: 0.1039
- val_mse: 0.0456

Epoch 42/50

46/46 [=====] - 16s 337ms/step - loss: 0.2268 - tp:
1329.0000 - fp: 129.0000 - tn: 1329.0000 - fn: 129.0000 - accuracy: 0.9115 -
precision: 0.9115 - auc: 0.9684 - mae: 0.1351 - mse: 0.0668 - val_loss: 0.1439 -
val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 -
val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9891 - val_mae: 0.0973
- val_mse: 0.0394

Epoch 43/50

46/46 [=====] - 16s 337ms/step - loss: 0.2318 - tp:
1334.0000 - fp: 124.0000 - tn: 1334.0000 - fn: 124.0000 - accuracy: 0.9150 -
precision: 0.9150 - auc: 0.9670 - mae: 0.1385 - mse: 0.0669 - val_loss: 0.1750 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9779 - val_mae: 0.0967
- val_mse: 0.0457

Epoch 44/50

46/46 [=====] - 15s 336ms/step - loss: 0.2322 - tp:
1335.0000 - fp: 123.0000 - tn: 1335.0000 - fn: 123.0000 - accuracy: 0.9156 -
precision: 0.9156 - auc: 0.9665 - mae: 0.1366 - mse: 0.0656 - val_loss: 0.1869 -
val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 -
val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9798 - val_mae: 0.1221
- val_mse: 0.0534

Epoch 45/50

46/46 [=====] - 15s 336ms/step - loss: 0.2472 - tp:
1304.0000 - fp: 154.0000 - tn: 1304.0000 - fn: 154.0000 - accuracy: 0.8944 -
precision: 0.8944 - auc: 0.9624 - mae: 0.1496 - mse: 0.0745 - val_loss: 0.2966 -
val_tp: 111.0000 - val_fp: 15.0000 - val_tn: 111.0000 - val_fn: 15.0000 -
val_accuracy: 0.8810 - val_precision: 0.8810 - val_auc: 0.9577 - val_mae: 0.1384
- val_mse: 0.0907

Epoch 46/50

46/46 [=====] - 16s 356ms/step - loss: 0.2202 - tp:
1330.0000 - fp: 128.0000 - tn: 1330.0000 - fn: 128.0000 - accuracy: 0.9122 -
precision: 0.9122 - auc: 0.9703 - mae: 0.1247 - mse: 0.0659 - val_loss: 0.1660 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9796 - val_mae: 0.0994
- val_mse: 0.0451

Epoch 47/50

46/46 [=====] - 15s 337ms/step - loss: 0.1867 - tp:
1355.0000 - fp: 103.0000 - tn: 1355.0000 - fn: 103.0000 - accuracy: 0.9294 -
precision: 0.9294 - auc: 0.9779 - mae: 0.1045 - mse: 0.0534 - val_loss: 0.1484 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9862 - val_mae: 0.0911
- val_mse: 0.0419

Epoch 48/50

46/46 [=====] - 15s 337ms/step - loss: 0.1854 - tp:
1360.0000 - fp: 98.0000 - tn: 1360.0000 - fn: 98.0000 - accuracy: 0.9328 -

```
precision: 0.9328 - auc: 0.9782 - mae: 0.1090 - mse: 0.0534 - val_loss: 0.1619 -  
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -  
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9840 - val_mae: 0.0862  
- val_mse: 0.0451
```

Epoch 49/50


```
46/46 [=====] - 15s 337ms/step - loss: 0.2046 - tp:  
1342.0000 - fp: 116.0000 - tn: 1342.0000 - fn: 116.0000 - accuracy: 0.9204 -  
precision: 0.9204 - auc: 0.9740 - mae: 0.1198 - mse: 0.0599 - val_loss: 0.1784 -  
val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 -  
val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9792 - val_mae: 0.1056  
- val_mse: 0.0495
```


Epoch 50/50

```
46/46 [=====] - 15s 336ms/step - loss: 0.1822 - tp:  
1354.0000 - fp: 104.0000 - tn: 1354.0000 - fn: 104.0000 - accuracy: 0.9287 -  
precision: 0.9287 - auc: 0.9797 - mae: 0.1066 - mse: 0.0543 - val_loss: 0.2075 -  
val_tp: 114.0000 - val_fp: 12.0000 - val_tn: 114.0000 - val_fn: 12.0000 -  
val_accuracy: 0.9048 - val_precision: 0.9048 - val_auc: 0.9716 - val_mae: 0.1031  
- val_mse: 0.0596
```

1.1.10 6. Display History 2

```
[18]: print_history(history2)
```





1.1.11 7. Evaluate the model 2

```
[19]: print_model_evaluation(model2);
```

```
714/714 [=====] - 3s 4ms/step - loss: 0.2949 - tp:
650.0000 - fp: 64.0000 - tn: 650.0000 - fn: 64.0000 - accuracy: 0.9104 -
precision: 0.9104 - auc: 0.9596 - mae: 0.1102 - mse: 0.0753
loss : 0.295
tp : 650.0
```

```
fp : 64.0
tn : 650.0
fn : 64.0
accuracy : 0.91
precision : 0.91
auc : 0.96
mae : 0.11
mse : 0.075
```

1.1.12 8. Predict with model 2

```
[20]: predict_and_print_roc(model2);
```

Label Predictions:

```
[1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1,
1, 0, 1, 0, 0, 0]
```

Real Labels:

```
[1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1,
1, 0, 1, 0, 0, 0]
```

Confusion Matrix:

```
[[13  2]
 [ 0 17]]
```

Accuracy: 0.94

Apple_unhealthy



Apple_healthy



Apple_healthy




Apple_healthy



Apple_healthy



Apple_unhealthy



Apple_healthy




Apple_unhealthy



Apple_unhealthy





1.1.13 3. Create and compile model 3

This model has like the previous model three convolution layers. It starts with a small number of neurons and gradually increases its neuron amount. In between every convolution layer, a max pooling layer manipulates the values. Additionally, 20% of random values are dropped at the end of the hidden layers. The result from the hidden layers is then flattened with three dense layers that again gradually become smaller.

```
[21]: model3 = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(128, (3, 3),
        padding='same', activation='relu', use_bias=False),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Flatten(),
```

```

tf.keras.layers.Dense(16, activation='relu', use_bias=False),
tf.keras.layers.Dense(8, activation='relu', use_bias=False),
tf.keras.layers.Dense(num_classes, activation='softmax')
])

```

```

[22]: opt = tf.keras.optimizers.Adam(lr=LEARNING_RATE)
model3.compile(optimizer=opt,
               loss='categorical_crossentropy',
               metrics=METRICS)

```

1.1.14 4. Display model structure 3

```

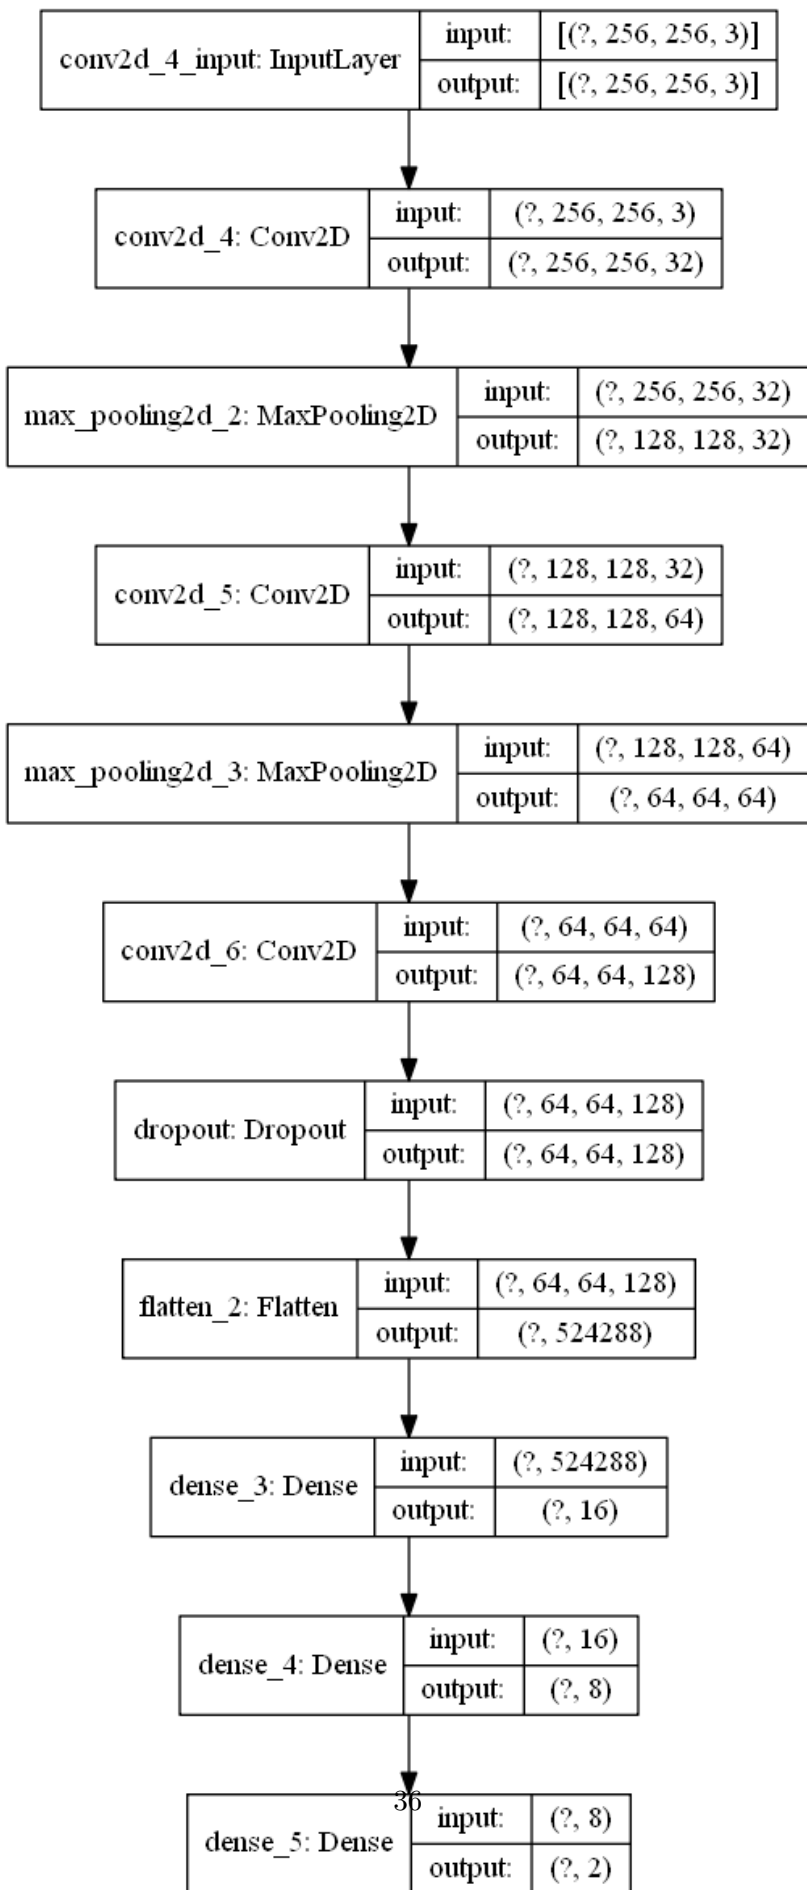
[23]: model3.summary()
tf.keras.utils.plot_model(model3, show_shapes=True)

```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|--------------------------------|----------------------|---------|
| conv2d_4 (Conv2D) | (None, 256, 256, 32) | 864 |
| max_pooling2d_2 (MaxPooling2D) | (None, 128, 128, 32) | 0 |
| conv2d_5 (Conv2D) | (None, 128, 128, 64) | 18432 |
| max_pooling2d_3 (MaxPooling2D) | (None, 64, 64, 64) | 0 |
| conv2d_6 (Conv2D) | (None, 64, 64, 128) | 73728 |
| dropout (Dropout) | (None, 64, 64, 128) | 0 |
| flatten_2 (Flatten) | (None, 524288) | 0 |
| dense_3 (Dense) | (None, 16) | 8388608 |
| dense_4 (Dense) | (None, 8) | 128 |
| dense_5 (Dense) | (None, 2) | 18 |
| Total params: 8,481,778 | | |
| Trainable params: 8,481,778 | | |
| Non-trainable params: 0 | | |

[23]:



1.1.15 5. Training model 3

```
[24]: history3 = train_model(model3, "Model_3")
```

Epoch 1/50

46/46 [=====] - 16s 356ms/step - loss: 0.7928 - tp: 1475.0000 - fp: 697.0000 - tn: 1475.0000 - fn: 697.0000 - accuracy: 0.6791 - precision: 0.6791 - auc: 0.7693 - mae: 0.3551 - mse: 0.1954 - val_loss: 0.5186 - val_tp: 106.0000 - val_fp: 20.0000 - val_tn: 106.0000 - val_fn: 20.0000 - val_accuracy: 0.8413 - val_precision: 0.8413 - val_auc: 0.8693 - val_mae: 0.3842 - val_mse: 0.1738

Epoch 2/50

46/46 [=====] - 15s 336ms/step - loss: 0.5012 - tp: 1149.0000 - fp: 309.0000 - tn: 1149.0000 - fn: 309.0000 - accuracy: 0.7881 - precision: 0.7881 - auc: 0.8445 - mae: 0.3200 - mse: 0.1626 - val_loss: 0.2349 - val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 - val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9697 - val_mae: 0.1596 - val_mse: 0.0633

Epoch 3/50

46/46 [=====] - 15s 336ms/step - loss: 0.3193 - tp: 1268.0000 - fp: 190.0000 - tn: 1268.0000 - fn: 190.0000 - accuracy: 0.8697 - precision: 0.8697 - auc: 0.9383 - mae: 0.1915 - mse: 0.0964 - val_loss: 0.2989 - val_tp: 111.0000 - val_fp: 15.0000 - val_tn: 111.0000 - val_fn: 15.0000 - val_accuracy: 0.8810 - val_precision: 0.8810 - val_auc: 0.9448 - val_mae: 0.1800 - val_mse: 0.0942

Epoch 4/50

46/46 [=====] - 16s 337ms/step - loss: 0.2667 - tp: 1300.0000 - fp: 158.0000 - tn: 1300.0000 - fn: 158.0000 - accuracy: 0.8916 - precision: 0.8916 - auc: 0.9568 - mae: 0.1661 - mse: 0.0807 - val_loss: 0.2337 - val_tp: 115.0000 - val_fp: 11.0000 - val_tn: 115.0000 - val_fn: 11.0000 - val_accuracy: 0.9127 - val_precision: 0.9127 - val_auc: 0.9677 - val_mae: 0.1440 - val_mse: 0.0648

Epoch 5/50

46/46 [=====] - 15s 336ms/step - loss: 0.2768 - tp: 1287.0000 - fp: 171.0000 - tn: 1287.0000 - fn: 171.0000 - accuracy: 0.8827 - precision: 0.8827 - auc: 0.9531 - mae: 0.1712 - mse: 0.0847 - val_loss: 0.2156 - val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 - val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9732 - val_mae: 0.1203 - val_mse: 0.0645

Epoch 6/50

46/46 [=====] - 15s 336ms/step - loss: 0.2313 - tp: 1325.0000 - fp: 133.0000 - tn: 1325.0000 - fn: 133.0000 - accuracy: 0.9088 - precision: 0.9088 - auc: 0.9678 - mae: 0.1384 - mse: 0.0679 - val_loss: 0.1876 - val_tp: 113.0000 - val_fp: 13.0000 - val_tn: 113.0000 - val_fn: 13.0000 - val_accuracy: 0.8968 - val_precision: 0.8968 - val_auc: 0.9792 - val_mae: 0.1202

```

- val_mse: 0.0571
Epoch 7/50
46/46 [=====] - 15s 336ms/step - loss: 0.2256 - tp:
1316.0000 - fp: 142.0000 - tn: 1316.0000 - fn: 142.0000 - accuracy: 0.9026 -
precision: 0.9026 - auc: 0.9693 - mae: 0.1342 - mse: 0.0704 - val_loss: 0.1771 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9795 - val_mae: 0.1113
- val_mse: 0.0455
Epoch 8/50
46/46 [=====] - 15s 337ms/step - loss: 0.2123 - tp:
1344.0000 - fp: 114.0000 - tn: 1344.0000 - fn: 114.0000 - accuracy: 0.9218 -
precision: 0.9218 - auc: 0.9723 - mae: 0.1265 - mse: 0.0621 - val_loss: 0.1875 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9774 - val_mae: 0.1129
- val_mse: 0.0550
Epoch 9/50
46/46 [=====] - 15s 336ms/step - loss: 0.2096 - tp:
1339.0000 - fp: 119.0000 - tn: 1339.0000 - fn: 119.0000 - accuracy: 0.9184 -
precision: 0.9184 - auc: 0.9730 - mae: 0.1252 - mse: 0.0613 - val_loss: 0.1542 -
val_tp: 123.0000 - val_fp: 3.0000 - val_tn: 123.0000 - val_fn: 3.0000 -
val_accuracy: 0.9762 - val_precision: 0.9762 - val_auc: 0.9879 - val_mae: 0.1076
- val_mse: 0.0347
Epoch 10/50
46/46 [=====] - 15s 337ms/step - loss: 0.2674 - tp:
1297.0000 - fp: 161.0000 - tn: 1297.0000 - fn: 161.0000 - accuracy: 0.8896 -
precision: 0.8896 - auc: 0.9562 - mae: 0.1605 - mse: 0.0809 - val_loss: 0.1716 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9799 - val_mae: 0.1057
- val_mse: 0.0419
Epoch 11/50
46/46 [=====] - 15s 336ms/step - loss: 0.2047 - tp:
1340.0000 - fp: 118.0000 - tn: 1340.0000 - fn: 118.0000 - accuracy: 0.9191 -
precision: 0.9191 - auc: 0.9744 - mae: 0.1249 - mse: 0.0600 - val_loss: 0.1713 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9832 - val_mae: 0.1120
- val_mse: 0.0422
Epoch 12/50
46/46 [=====] - 15s 336ms/step - loss: 0.1853 - tp:
1356.0000 - fp: 102.0000 - tn: 1356.0000 - fn: 102.0000 - accuracy: 0.9300 -
precision: 0.9300 - auc: 0.9786 - mae: 0.1089 - mse: 0.0532 - val_loss: 0.1316 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9921 - val_mae: 0.0964
- val_mse: 0.0365
Epoch 13/50
46/46 [=====] - 15s 336ms/step - loss: 0.1788 - tp:
1354.0000 - fp: 104.0000 - tn: 1354.0000 - fn: 104.0000 - accuracy: 0.9287 -
precision: 0.9287 - auc: 0.9800 - mae: 0.1032 - mse: 0.0521 - val_loss: 0.1179 -
val_tp: 121.0000 - val_fp: 5.0000 - val_tn: 121.0000 - val_fn: 5.0000 -

```

val_accuracy: 0.9603 - val_precision: 0.9603 - val_auc: 0.9949 - val_mae: 0.0912
- val_mse: 0.0318

Epoch 14/50

46/46 [=====] - 15s 337ms/step - loss: 0.1901 - tp:
1352.0000 - fp: 106.0000 - tn: 1352.0000 - fn: 106.0000 - accuracy: 0.9273 -
precision: 0.9273 - auc: 0.9775 - mae: 0.1081 - mse: 0.0554 - val_loss: 0.2081 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9735 - val_mae: 0.1125
- val_mse: 0.0558

Epoch 15/50

46/46 [=====] - 16s 339ms/step - loss: 0.1688 - tp:
1362.0000 - fp: 96.0000 - tn: 1362.0000 - fn: 96.0000 - accuracy: 0.9342 -
precision: 0.9342 - auc: 0.9823 - mae: 0.1017 - mse: 0.0487 - val_loss: 0.1201 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9906 - val_mae: 0.0758
- val_mse: 0.0286

Epoch 16/50

46/46 [=====] - 15s 336ms/step - loss: 0.1528 - tp:
1370.0000 - fp: 88.0000 - tn: 1370.0000 - fn: 88.0000 - accuracy: 0.9396 -
precision: 0.9396 - auc: 0.9857 - mae: 0.0922 - mse: 0.0446 - val_loss: 0.1256 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9918 - val_mae: 0.0855
- val_mse: 0.0347

Epoch 17/50

46/46 [=====] - 15s 336ms/step - loss: 0.1270 - tp:
1387.0000 - fp: 71.0000 - tn: 1387.0000 - fn: 71.0000 - accuracy: 0.9513 -
precision: 0.9513 - auc: 0.9896 - mae: 0.0769 - mse: 0.0357 - val_loss: 0.1224 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9912 - val_mae: 0.0802
- val_mse: 0.0376

Epoch 18/50

46/46 [=====] - 15s 336ms/step - loss: 0.1452 - tp:
1375.0000 - fp: 83.0000 - tn: 1375.0000 - fn: 83.0000 - accuracy: 0.9431 -
precision: 0.9431 - auc: 0.9873 - mae: 0.0802 - mse: 0.0429 - val_loss: 0.1627 -
val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 -
val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9852 - val_mae: 0.1114
- val_mse: 0.0468

Epoch 19/50

46/46 [=====] - 15s 336ms/step - loss: 0.1505 - tp:
1368.0000 - fp: 90.0000 - tn: 1368.0000 - fn: 90.0000 - accuracy: 0.9383 -
precision: 0.9383 - auc: 0.9856 - mae: 0.0901 - mse: 0.0442 - val_loss: 0.1131 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9921 - val_mae: 0.0646
- val_mse: 0.0336

Epoch 20/50

46/46 [=====] - 15s 336ms/step - loss: 0.1284 - tp:
1383.0000 - fp: 75.0000 - tn: 1383.0000 - fn: 75.0000 - accuracy: 0.9486 -
precision: 0.9486 - auc: 0.9893 - mae: 0.0792 - mse: 0.0373 - val_loss: 0.1152 -

val_tp: 121.0000 - val_fp: 5.0000 - val_tn: 121.0000 - val_fn: 5.0000 -
val_accuracy: 0.9603 - val_precision: 0.9603 - val_auc: 0.9917 - val_mae: 0.0644
- val_mse: 0.0326

Epoch 21/50

46/46 [=====] - 15s 336ms/step - loss: 0.1126 - tp:
1397.0000 - fp: 61.0000 - tn: 1397.0000 - fn: 61.0000 - accuracy: 0.9582 -
precision: 0.9582 - auc: 0.9918 - mae: 0.0689 - mse: 0.0323 - val_loss: 0.1342 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9880 - val_mae: 0.0655
- val_mse: 0.0356

Epoch 22/50

46/46 [=====] - 15s 336ms/step - loss: 0.1593 - tp:
1366.0000 - fp: 92.0000 - tn: 1366.0000 - fn: 92.0000 - accuracy: 0.9369 -
precision: 0.9369 - auc: 0.9838 - mae: 0.0911 - mse: 0.0466 - val_loss: 0.1058 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9936 - val_mae: 0.0614
- val_mse: 0.0316

Epoch 23/50

46/46 [=====] - 15s 337ms/step - loss: 0.1158 - tp:
1387.0000 - fp: 71.0000 - tn: 1387.0000 - fn: 71.0000 - accuracy: 0.9513 -
precision: 0.9513 - auc: 0.9914 - mae: 0.0672 - mse: 0.0350 - val_loss: 0.1229 -
val_tp: 121.0000 - val_fp: 5.0000 - val_tn: 121.0000 - val_fn: 5.0000 -
val_accuracy: 0.9603 - val_precision: 0.9603 - val_auc: 0.9905 - val_mae: 0.0673
- val_mse: 0.0367

Epoch 24/50

46/46 [=====] - 15s 336ms/step - loss: 0.0934 - tp:
1415.0000 - fp: 43.0000 - tn: 1415.0000 - fn: 43.0000 - accuracy: 0.9705 -
precision: 0.9705 - auc: 0.9941 - mae: 0.0578 - mse: 0.0256 - val_loss: 0.0988 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9890 - val_mae: 0.0390
- val_mse: 0.0241

Epoch 25/50

46/46 [=====] - 15s 336ms/step - loss: 0.1272 - tp:
1399.0000 - fp: 59.0000 - tn: 1399.0000 - fn: 59.0000 - accuracy: 0.9595 -
precision: 0.9595 - auc: 0.9896 - mae: 0.0678 - mse: 0.0358 - val_loss: 0.1041 -
val_tp: 121.0000 - val_fp: 5.0000 - val_tn: 121.0000 - val_fn: 5.0000 -
val_accuracy: 0.9603 - val_precision: 0.9603 - val_auc: 0.9926 - val_mae: 0.0563
- val_mse: 0.0272

Epoch 26/50

46/46 [=====] - 16s 338ms/step - loss: 0.1126 - tp:
1393.0000 - fp: 65.0000 - tn: 1393.0000 - fn: 65.0000 - accuracy: 0.9554 -
precision: 0.9554 - auc: 0.9913 - mae: 0.0642 - mse: 0.0323 - val_loss: 0.1520 -
val_tp: 117.0000 - val_fp: 9.0000 - val_tn: 117.0000 - val_fn: 9.0000 -
val_accuracy: 0.9286 - val_precision: 0.9286 - val_auc: 0.9879 - val_mae: 0.0745
- val_mse: 0.0461

Epoch 27/50

46/46 [=====] - 15s 336ms/step - loss: 0.1266 - tp:
1392.0000 - fp: 66.0000 - tn: 1392.0000 - fn: 66.0000 - accuracy: 0.9547 -

precision: 0.9547 - auc: 0.9895 - mae: 0.0716 - mse: 0.0360 - val_loss: 0.1181 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9866 - val_mae: 0.0499
- val_mse: 0.0303

Epoch 28/50

46/46 [=====] - 15s 336ms/step - loss: 0.1007 - tp:
1402.0000 - fp: 56.0000 - tn: 1402.0000 - fn: 56.0000 - accuracy: 0.9616 -
precision: 0.9616 - auc: 0.9938 - mae: 0.0566 - mse: 0.0291 - val_loss: 0.1010 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9940 - val_mae: 0.0473
- val_mse: 0.0274

Epoch 29/50

46/46 [=====] - 15s 336ms/step - loss: 0.0749 - tp:
1414.0000 - fp: 44.0000 - tn: 1414.0000 - fn: 44.0000 - accuracy: 0.9698 -
precision: 0.9698 - auc: 0.9966 - mae: 0.0437 - mse: 0.0221 - val_loss: 0.1248 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9931 - val_mae: 0.0536
- val_mse: 0.0360

Epoch 30/50

46/46 [=====] - 15s 336ms/step - loss: 0.1030 - tp:
1405.0000 - fp: 53.0000 - tn: 1405.0000 - fn: 53.0000 - accuracy: 0.9636 -
precision: 0.9636 - auc: 0.9922 - mae: 0.0534 - mse: 0.0278 - val_loss: 0.1238 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9893 - val_mae: 0.0550
- val_mse: 0.0307

Epoch 31/50

46/46 [=====] - 16s 339ms/step - loss: 0.0822 - tp:
1408.0000 - fp: 50.0000 - tn: 1408.0000 - fn: 50.0000 - accuracy: 0.9657 -
precision: 0.9657 - auc: 0.9955 - mae: 0.0508 - mse: 0.0240 - val_loss: 0.1061 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9934 - val_mae: 0.0470
- val_mse: 0.0251

Epoch 32/50

46/46 [=====] - 16s 338ms/step - loss: 0.0967 - tp:
1405.0000 - fp: 53.0000 - tn: 1405.0000 - fn: 53.0000 - accuracy: 0.9636 -
precision: 0.9636 - auc: 0.9938 - mae: 0.0527 - mse: 0.0278 - val_loss: 0.1130 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9932 - val_mae: 0.0567
- val_mse: 0.0350

Epoch 33/50

46/46 [=====] - 16s 340ms/step - loss: 0.0923 - tp:
1411.0000 - fp: 47.0000 - tn: 1411.0000 - fn: 47.0000 - accuracy: 0.9678 -
precision: 0.9678 - auc: 0.9933 - mae: 0.0496 - mse: 0.0246 - val_loss: 0.1569 -
val_tp: 119.0000 - val_fp: 7.0000 - val_tn: 119.0000 - val_fn: 7.0000 -
val_accuracy: 0.9444 - val_precision: 0.9444 - val_auc: 0.9858 - val_mae: 0.0531
- val_mse: 0.0385

Epoch 34/50

46/46 [=====] - 15s 336ms/step - loss: 0.0551 - tp:

1431.0000 - fp: 27.0000 - tn: 1431.0000 - fn: 27.0000 - accuracy: 0.9815 -
precision: 0.9815 - auc: 0.9976 - mae: 0.0339 - mse: 0.0153 - val_loss: 0.1376 -
val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 -
val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9865 - val_mae: 0.0542
- val_mse: 0.0354

Epoch 35/50

46/46 [=====] - 15s 336ms/step - loss: 0.0533 - tp:
1426.0000 - fp: 32.0000 - tn: 1426.0000 - fn: 32.0000 - accuracy: 0.9781 -
precision: 0.9781 - auc: 0.9983 - mae: 0.0316 - mse: 0.0158 - val_loss: 0.1557 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9849 - val_mae: 0.0468
- val_mse: 0.0338

Epoch 36/50

46/46 [=====] - 15s 336ms/step - loss: 0.1060 - tp:
1398.0000 - fp: 60.0000 - tn: 1398.0000 - fn: 60.0000 - accuracy: 0.9588 -
precision: 0.9588 - auc: 0.9930 - mae: 0.0546 - mse: 0.0311 - val_loss: 0.1469 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9852 - val_mae: 0.0502
- val_mse: 0.0298

Epoch 37/50

46/46 [=====] - 15s 336ms/step - loss: 0.0485 - tp:
1432.0000 - fp: 26.0000 - tn: 1432.0000 - fn: 26.0000 - accuracy: 0.9822 -
precision: 0.9822 - auc: 0.9989 - mae: 0.0348 - mse: 0.0137 - val_loss: 0.1379 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9880 - val_mae: 0.0348
- val_mse: 0.0282

Epoch 38/50

46/46 [=====] - 15s 336ms/step - loss: 0.0742 - tp:
1417.0000 - fp: 41.0000 - tn: 1417.0000 - fn: 41.0000 - accuracy: 0.9719 -
precision: 0.9719 - auc: 0.9962 - mae: 0.0389 - mse: 0.0215 - val_loss: 0.0659 -
val_tp: 123.0000 - val_fp: 3.0000 - val_tn: 123.0000 - val_fn: 3.0000 -
val_accuracy: 0.9762 - val_precision: 0.9762 - val_auc: 0.9977 - val_mae: 0.0455
- val_mse: 0.0205

Epoch 39/50

46/46 [=====] - 15s 336ms/step - loss: 0.0947 - tp:
1405.0000 - fp: 53.0000 - tn: 1405.0000 - fn: 53.0000 - accuracy: 0.9636 -
precision: 0.9636 - auc: 0.9941 - mae: 0.0540 - mse: 0.0271 - val_loss: 0.1701 -
val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 -
val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9790 - val_mae: 0.0472
- val_mse: 0.0318

Epoch 40/50

46/46 [=====] - 15s 336ms/step - loss: 0.0858 - tp:
1408.0000 - fp: 50.0000 - tn: 1408.0000 - fn: 50.0000 - accuracy: 0.9657 -
precision: 0.9657 - auc: 0.9956 - mae: 0.0457 - mse: 0.0253 - val_loss: 0.3218 -
val_tp: 116.0000 - val_fp: 10.0000 - val_tn: 116.0000 - val_fn: 10.0000 -
val_accuracy: 0.9206 - val_precision: 0.9206 - val_auc: 0.9604 - val_mae: 0.0901
- val_mse: 0.0677

Epoch 41/50

46/46 [=====] - 16s 338ms/step - loss: 0.1026 - tp: 1402.0000 - fp: 56.0000 - tn: 1402.0000 - fn: 56.0000 - accuracy: 0.9616 - precision: 0.9616 - auc: 0.9933 - mae: 0.0562 - mse: 0.0301 - val_loss: 0.0863 - val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 - val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9955 - val_mae: 0.0472 - val_mse: 0.0252

Epoch 42/50

46/46 [=====] - 15s 337ms/step - loss: 0.0572 - tp: 1427.0000 - fp: 31.0000 - tn: 1427.0000 - fn: 31.0000 - accuracy: 0.9787 - precision: 0.9787 - auc: 0.9980 - mae: 0.0352 - mse: 0.0170 - val_loss: 0.1763 - val_tp: 121.0000 - val_fp: 5.0000 - val_tn: 121.0000 - val_fn: 5.0000 - val_accuracy: 0.9603 - val_precision: 0.9603 - val_auc: 0.9866 - val_mae: 0.0540 - val_mse: 0.0330

Epoch 43/50

46/46 [=====] - 15s 336ms/step - loss: 0.0924 - tp: 1410.0000 - fp: 48.0000 - tn: 1410.0000 - fn: 48.0000 - accuracy: 0.9671 - precision: 0.9671 - auc: 0.9941 - mae: 0.0474 - mse: 0.0251 - val_loss: 0.1121 - val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 - val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9892 - val_mae: 0.0393 - val_mse: 0.0230

Epoch 44/50

46/46 [=====] - 15s 336ms/step - loss: 0.0498 - tp: 1434.0000 - fp: 24.0000 - tn: 1434.0000 - fn: 24.0000 - accuracy: 0.9835 - precision: 0.9835 - auc: 0.9979 - mae: 0.0292 - mse: 0.0131 - val_loss: 0.0717 - val_tp: 123.0000 - val_fp: 3.0000 - val_tn: 123.0000 - val_fn: 3.0000 - val_accuracy: 0.9762 - val_precision: 0.9762 - val_auc: 0.9971 - val_mae: 0.0310 - val_mse: 0.0189

Epoch 45/50

46/46 [=====] - 15s 336ms/step - loss: 0.0385 - tp: 1438.0000 - fp: 20.0000 - tn: 1438.0000 - fn: 20.0000 - accuracy: 0.9863 - precision: 0.9863 - auc: 0.9991 - mae: 0.0242 - mse: 0.0105 - val_loss: 0.1966 - val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 - val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9819 - val_mae: 0.0607 - val_mse: 0.0454

Epoch 46/50

46/46 [=====] - 15s 336ms/step - loss: 0.0696 - tp: 1421.0000 - fp: 37.0000 - tn: 1421.0000 - fn: 37.0000 - accuracy: 0.9746 - precision: 0.9746 - auc: 0.9966 - mae: 0.0370 - mse: 0.0194 - val_loss: 0.1506 - val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 - val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9917 - val_mae: 0.0556 - val_mse: 0.0414

Epoch 47/50

46/46 [=====] - 15s 336ms/step - loss: 0.0550 - tp: 1425.0000 - fp: 33.0000 - tn: 1425.0000 - fn: 33.0000 - accuracy: 0.9774 - precision: 0.9774 - auc: 0.9977 - mae: 0.0338 - mse: 0.0164 - val_loss: 0.1395 - val_tp: 121.0000 - val_fp: 5.0000 - val_tn: 121.0000 - val_fn: 5.0000 - val_accuracy: 0.9603 - val_precision: 0.9603 - val_auc: 0.9795 - val_mae: 0.0439 - val_mse: 0.0306

Epoch 48/50

46/46 [=====] - 15s 336ms/step - loss: 0.0733 - tp: 1418.0000 - fp: 40.0000 - tn: 1418.0000 - fn: 40.0000 - accuracy: 0.9726 - precision: 0.9726 - auc: 0.9956 - mae: 0.0378 - mse: 0.0199 - val_loss: 0.1452 - val_tp: 122.0000 - val_fp: 4.0000 - val_tn: 122.0000 - val_fn: 4.0000 - val_accuracy: 0.9683 - val_precision: 0.9683 - val_auc: 0.9873 - val_mae: 0.0370 - val_mse: 0.0294

Epoch 49/50


46/46 [=====] - 15s 335ms/step - loss: 0.0418 - tp: 1439.0000 - fp: 19.0000 - tn: 1439.0000 - fn: 19.0000 - accuracy: 0.9870 - precision: 0.9870 - auc: 0.9988 - mae: 0.0238 - mse: 0.0109 - val_loss: 0.1659 - val_tp: 120.0000 - val_fp: 6.0000 - val_tn: 120.0000 - val_fn: 6.0000 - val_accuracy: 0.9524 - val_precision: 0.9524 - val_auc: 0.9839 - val_mae: 0.0575 - val_mse: 0.0413

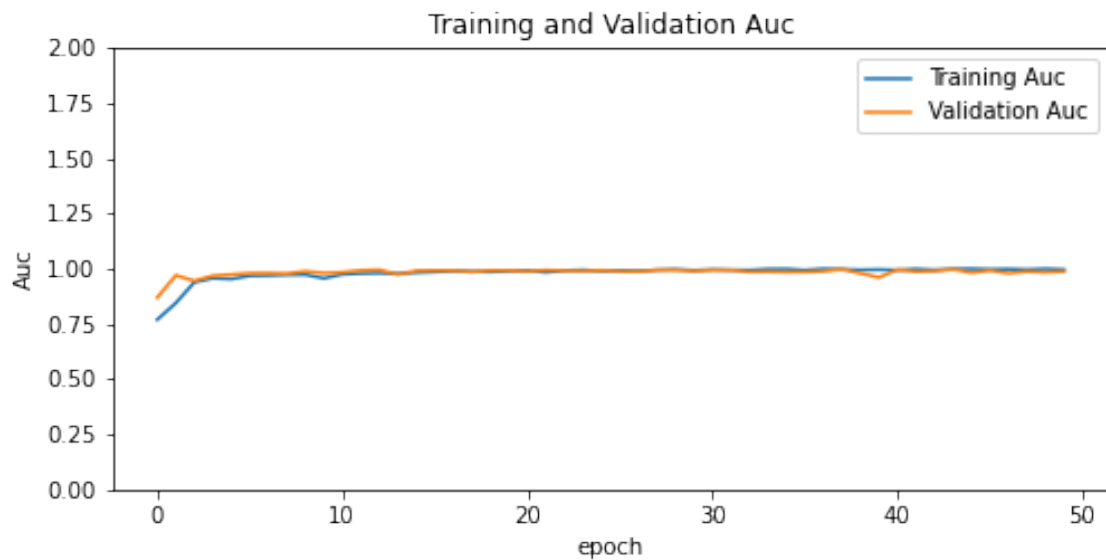

Epoch 50/50

46/46 [=====] - 15s 335ms/step - loss: 0.0684 - tp: 1424.0000 - fp: 34.0000 - tn: 1424.0000 - fn: 34.0000 - accuracy: 0.9767 - precision: 0.9767 - auc: 0.9957 - mae: 0.0343 - mse: 0.0186 - val_loss: 0.1548 - val_tp: 118.0000 - val_fp: 8.0000 - val_tn: 118.0000 - val_fn: 8.0000 - val_accuracy: 0.9365 - val_precision: 0.9365 - val_auc: 0.9874 - val_mae: 0.0773 - val_mse: 0.0470

1.1.16 6. Display History 3

```
[25]: print_history(history3)
```





1.1.17 7. Evaluate the model 3

```
[26]: print_model_evaluation(model3);
```

```
714/714 [=====] - 3s 4ms/step - loss: 0.1899 - tp:
669.0000 - fp: 45.0000 - tn: 669.0000 - fn: 45.0000 - accuracy: 0.9370 -
precision: 0.9370 - auc: 0.9823 - mae: 0.0709 - mse: 0.0514
loss : 0.19
tp : 669.0
```

```
fp : 45.0
tn : 669.0
fn : 45.0
accuracy : 0.937
precision : 0.937
auc : 0.982
mae : 0.071
mse : 0.051
```

1.1.18 8. Predict with model 3

```
[27]: predict_and_print_roc(model3);
```

Label Predictions:

```
[1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0,
0, 1, 1, 0, 0, 1]
```

Real Labels:

```
[1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0,
0, 1, 1, 0, 0, 1]
```

Confusion Matrix:

```
[[14  2]
 [ 0 16]]
```

Accuracy: 0.94

Apple_unhealthy



Apple_healthy



Apple_healthy



Apple_healthy



Apple_unhealthy



Apple_healthy



Apple_unhealthy




Apple_unhealthy



Apple_unhealthy





1.1.19 9. Save test model

```
[28]: model3.save(MODEL_NAME)
```

1.1.20 10. Load test model

```
[29]: loaded_model = tf.keras.models.load_model(MODEL_NAME)
```

1.1.21 11. Test loaded model on image:

```
[30]: healthyTestImagePath = "./plant_images/test/Apple__healthy/
      ↳2d7f4c07-f4f4-4589-aa82-49e05e5b253b__RS_HL_8056.jpg"

healthyTestImage = image.load_img(healthyTestImagePath)
plt.imshow(healthyTestImage)
healthyTestImage = (np.expand_dims(healthyTestImage,0))
predictions = loaded_model.predict(healthyTestImage)


print('Expected result: [1, 0] \n')
```



```
print('Result: ', predictions, '\n')
```

Expected result: [1, 0]


Result: [[1. 0.]]



```
[34]: unhealthyTestImagePath = "./plant_images/test/Apple__unhealthy/  
      ↳0ebea6f4-08e4-4380-86f8-34d854697e32___JR_FrgE.S 2877.jpg"  
unhealthyTestImage = image.load_img(unhealthyTestImagePath)  
plt.imshow(unhealthyTestImage)  
unhealthyTestImage = (np.expand_dims(unhealthyTestImage,0))  
predictions = loaded_model.predict(unhealthyTestImage)  
  
print('Expected result: [0, 1] \n')  
print('Result: ', predictions, '\n')
```

Expected result: [0, 1]

Result: [[0. 1.]]



1.2 Summary

All the models were trained with 50 epochs. The reason for this is that after this amount no special improvements could be observed anymore and overfitting started to become a problem (tested with 100 epochs before).

1.2.1 Model 1:

This is a very simple model that already works very well. Showing that a small amount of layers can already create a very good model.

1.2.2 Model 2:

Has more hidden and dense layers than model 1. Increasing the amount of neurons each layer, however, didn't seem to work as well as expected. Nonetheless, more layers created a better accuracy overall.

1.2.3 Model 3:

This model seems to perform the best, based on the false positive and false negative validation. Decreasing the size of the image and the neurons and leaving the amount of layers around 2-3 seemed to have the best effect.

[]: